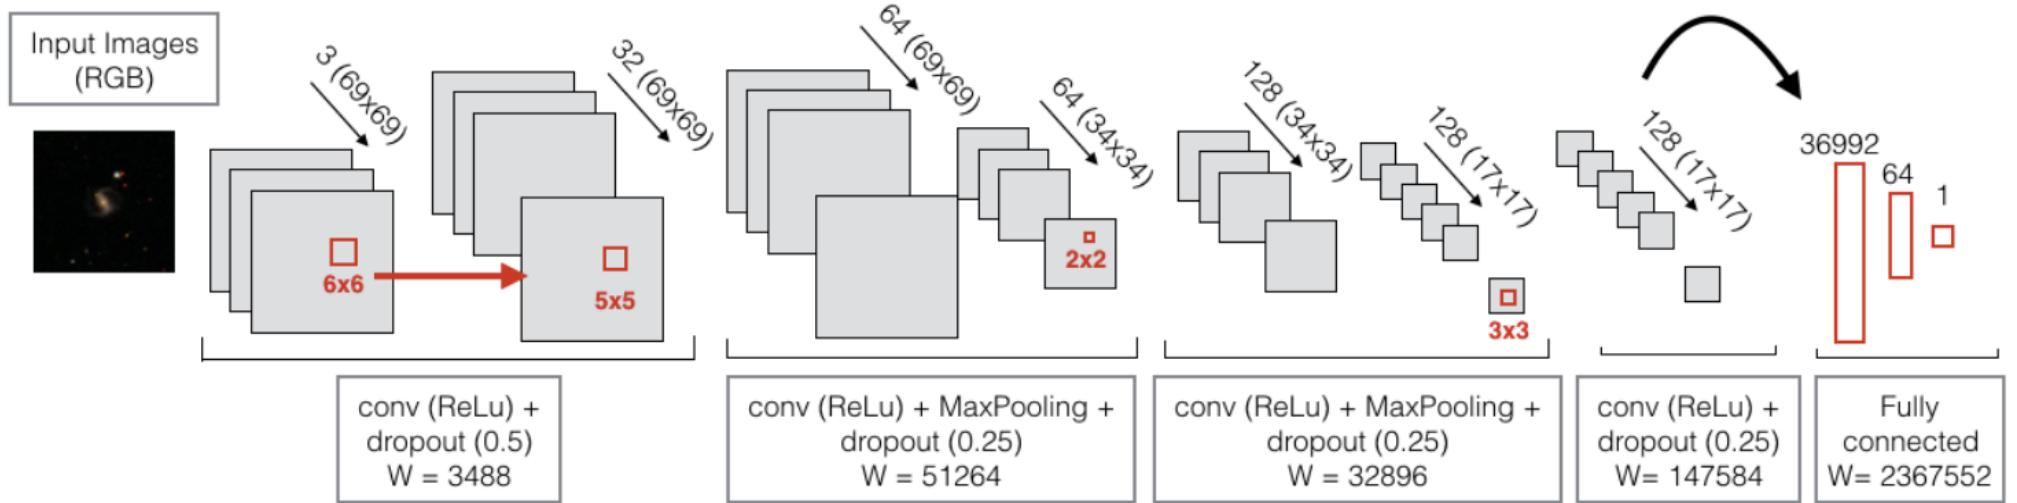


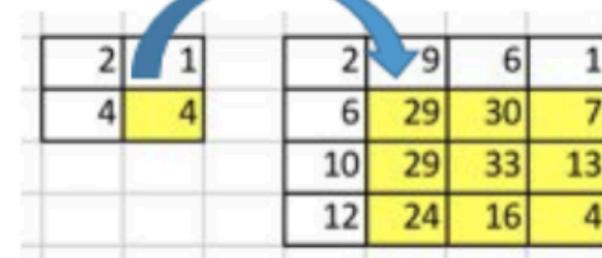
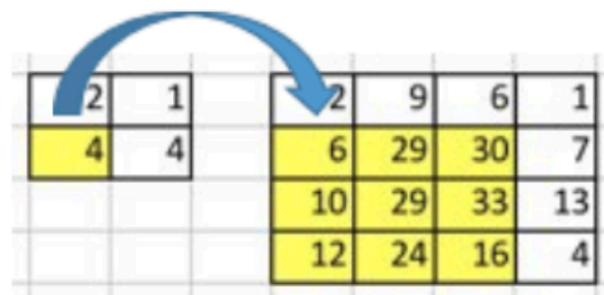
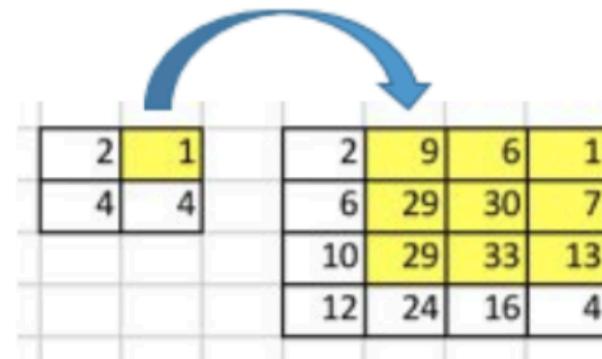
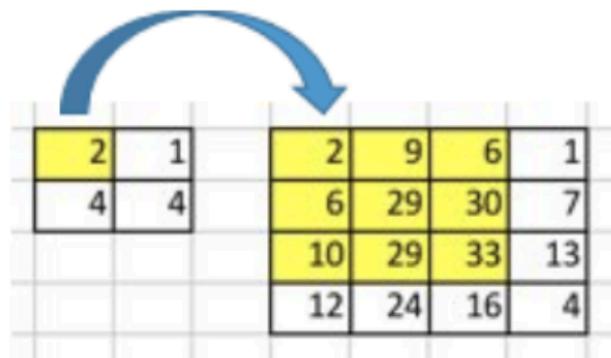
EXAMPLE OF VERY SIMPLE CNN



Dominguez-Sanchez+18

TRANSPOSED CONVOLUTION

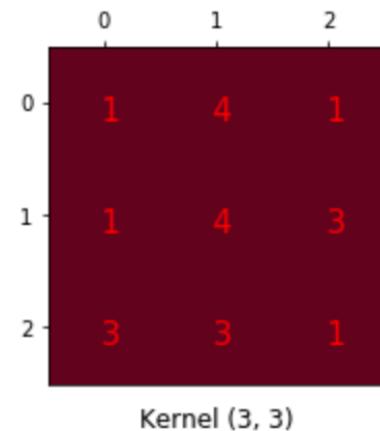
ALLOWS TO INCREASE THE SIZE



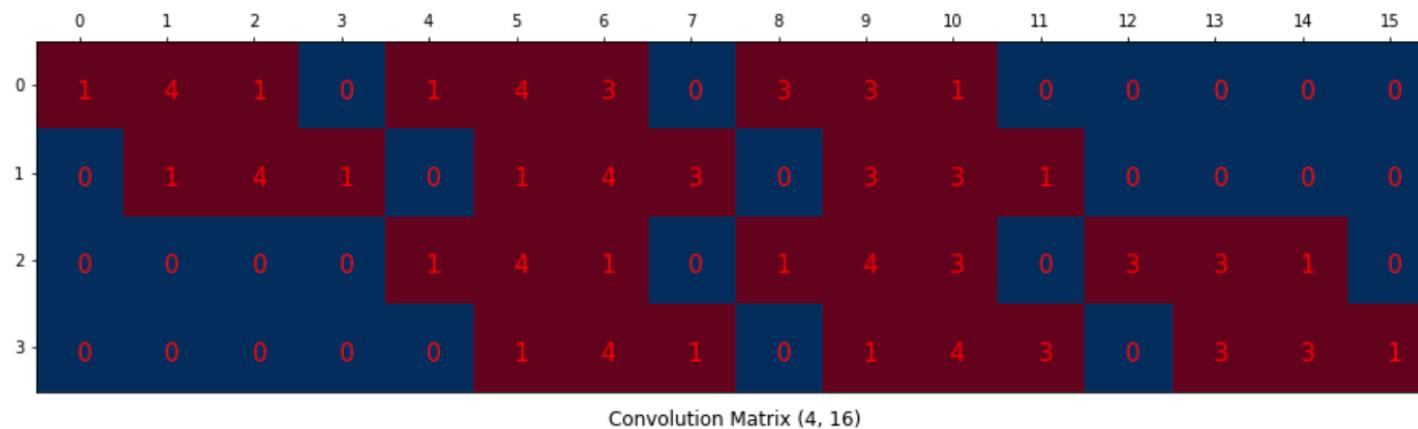
Going Backward of Convolution

EXAMPLE TAKEN FROM HERE

CONVOLUTION MATRIX

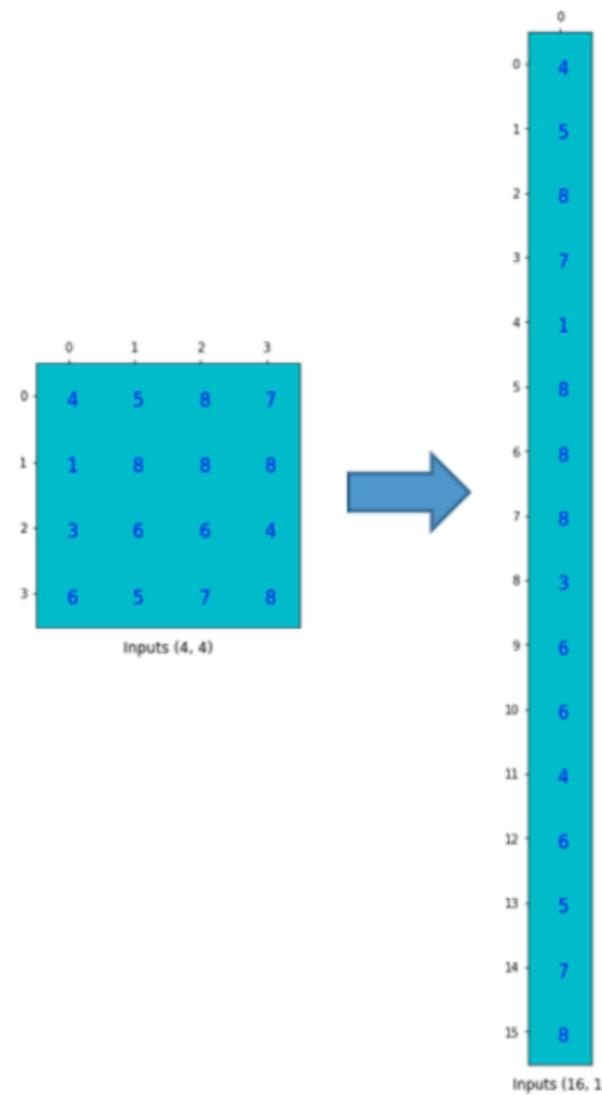


THE KERNEL CAN BE ARRANGED IN FORM OF A MATRIX:



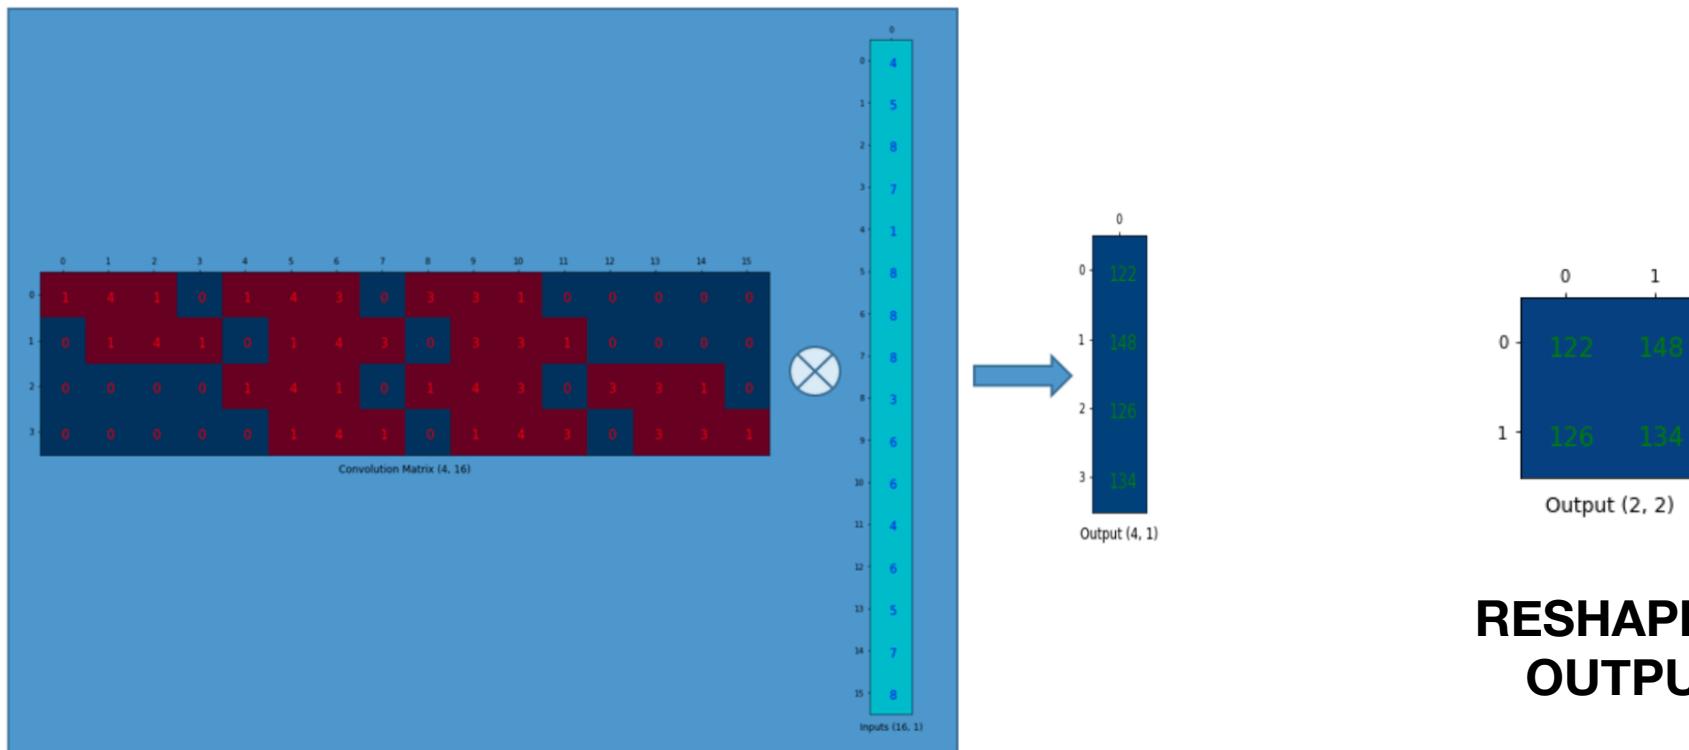
EXAMPLE TAKEN FROM HERE

THE INPUT IS FLATTENED INTO A COLUMN VECTOR

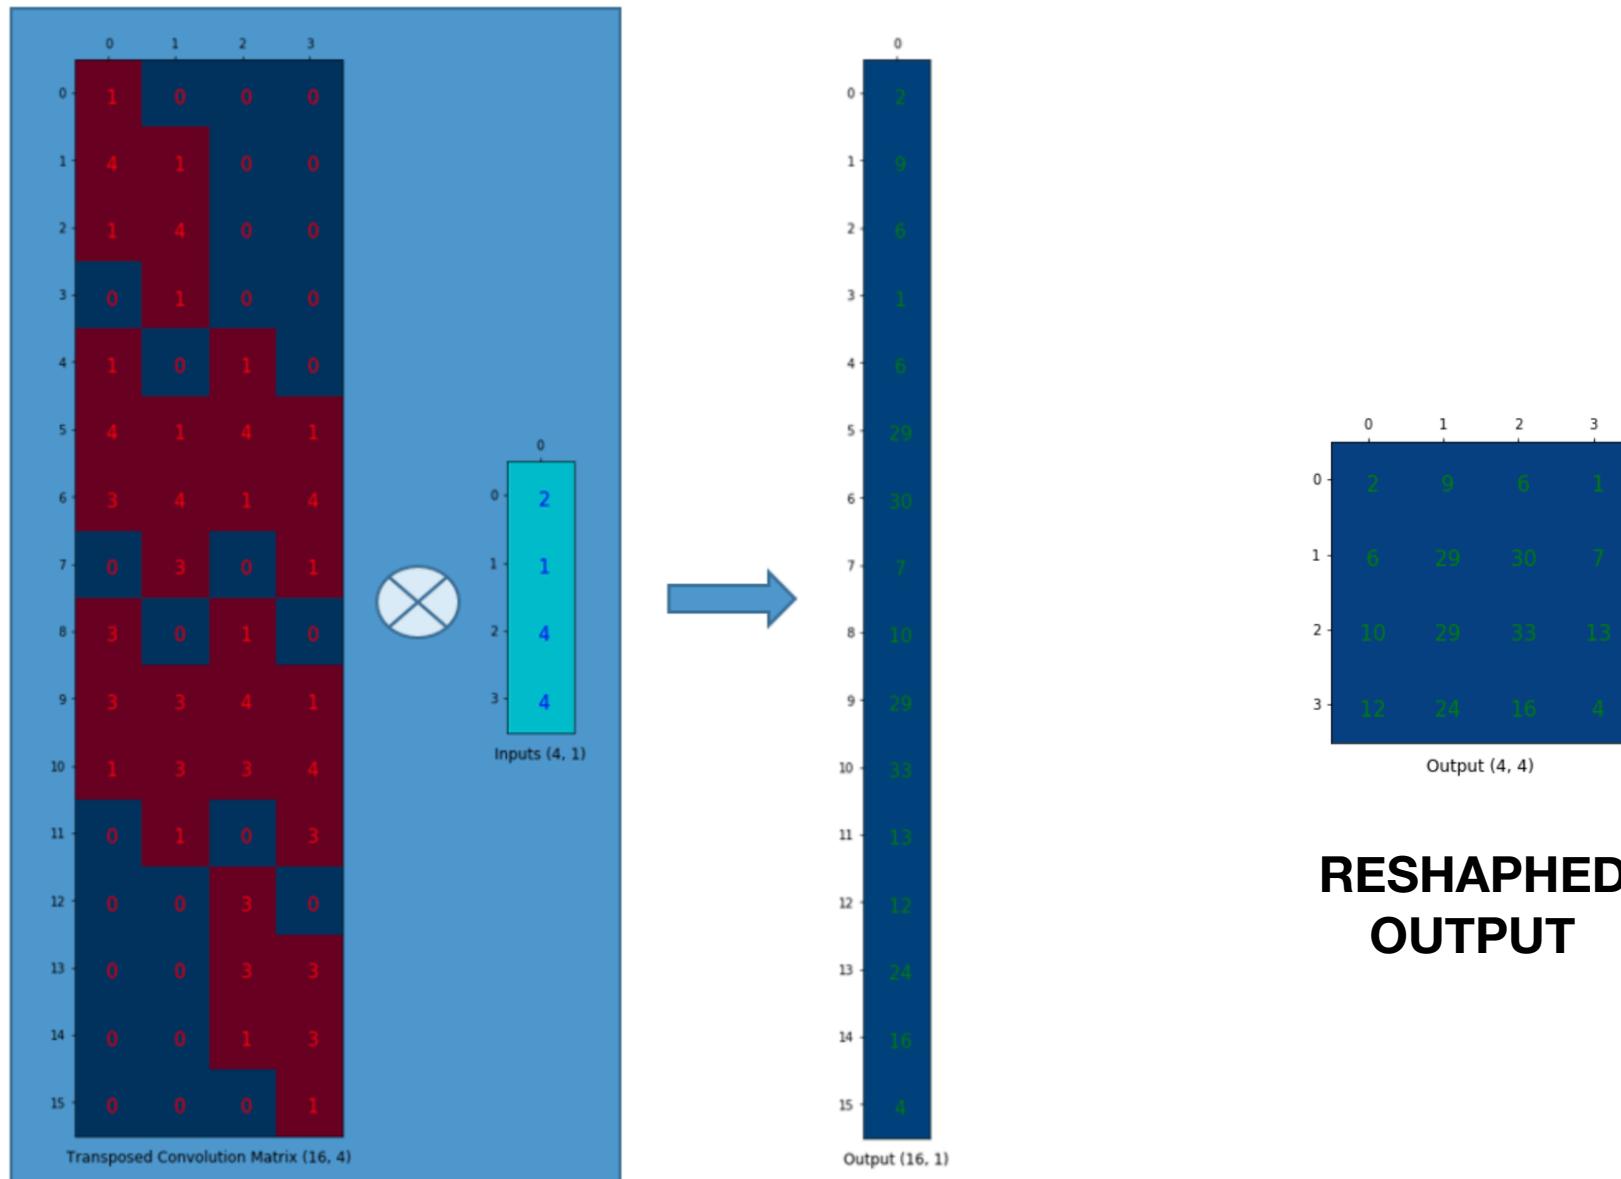


EXAMPLE TAKEN FROM HERE

THE CONVOLUTION IS TRANSFORMED INTO A PRODUCT OF MATRICES

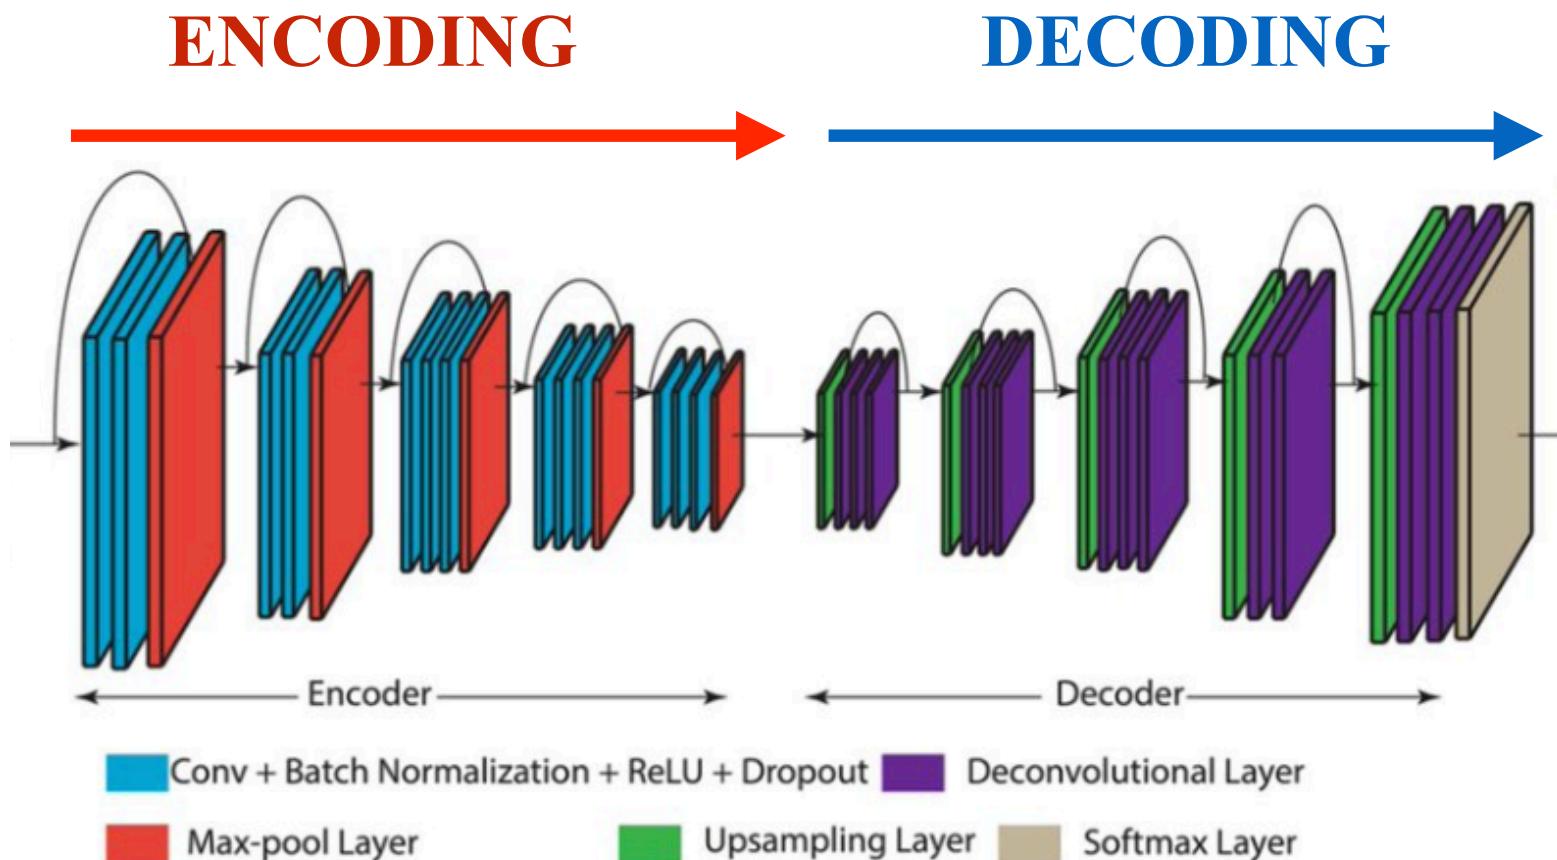


THE TRANSPOSED CONVOLUTION IS THE INVERSE OPERATION

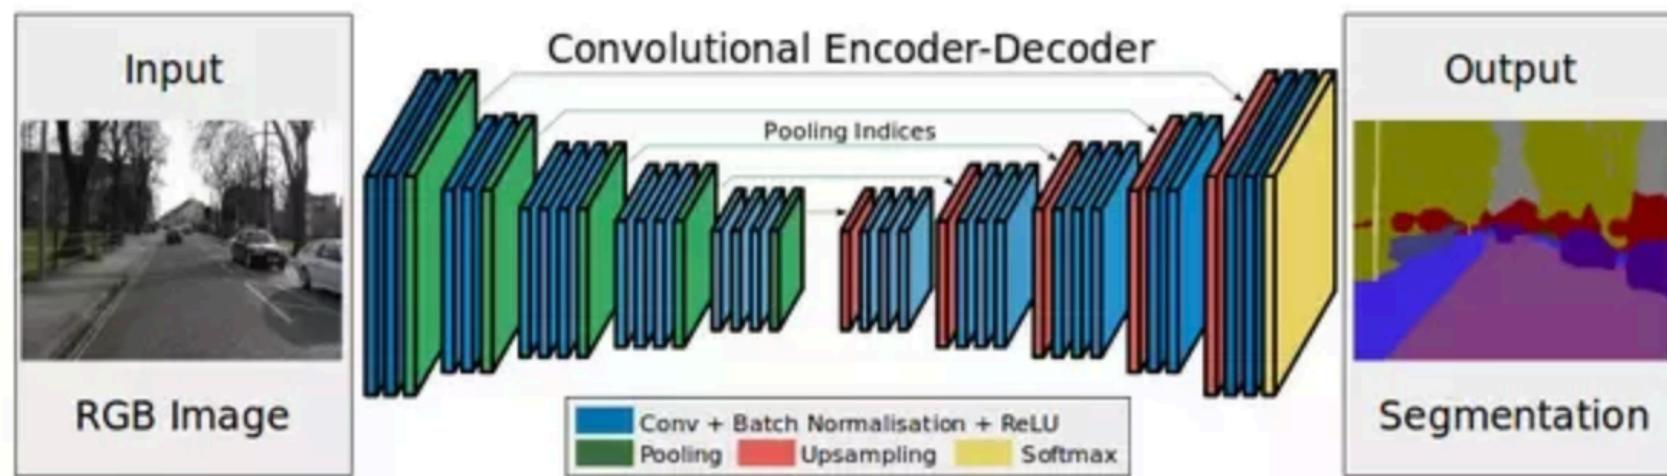


EXAMPLE TAKEN FROM HERE

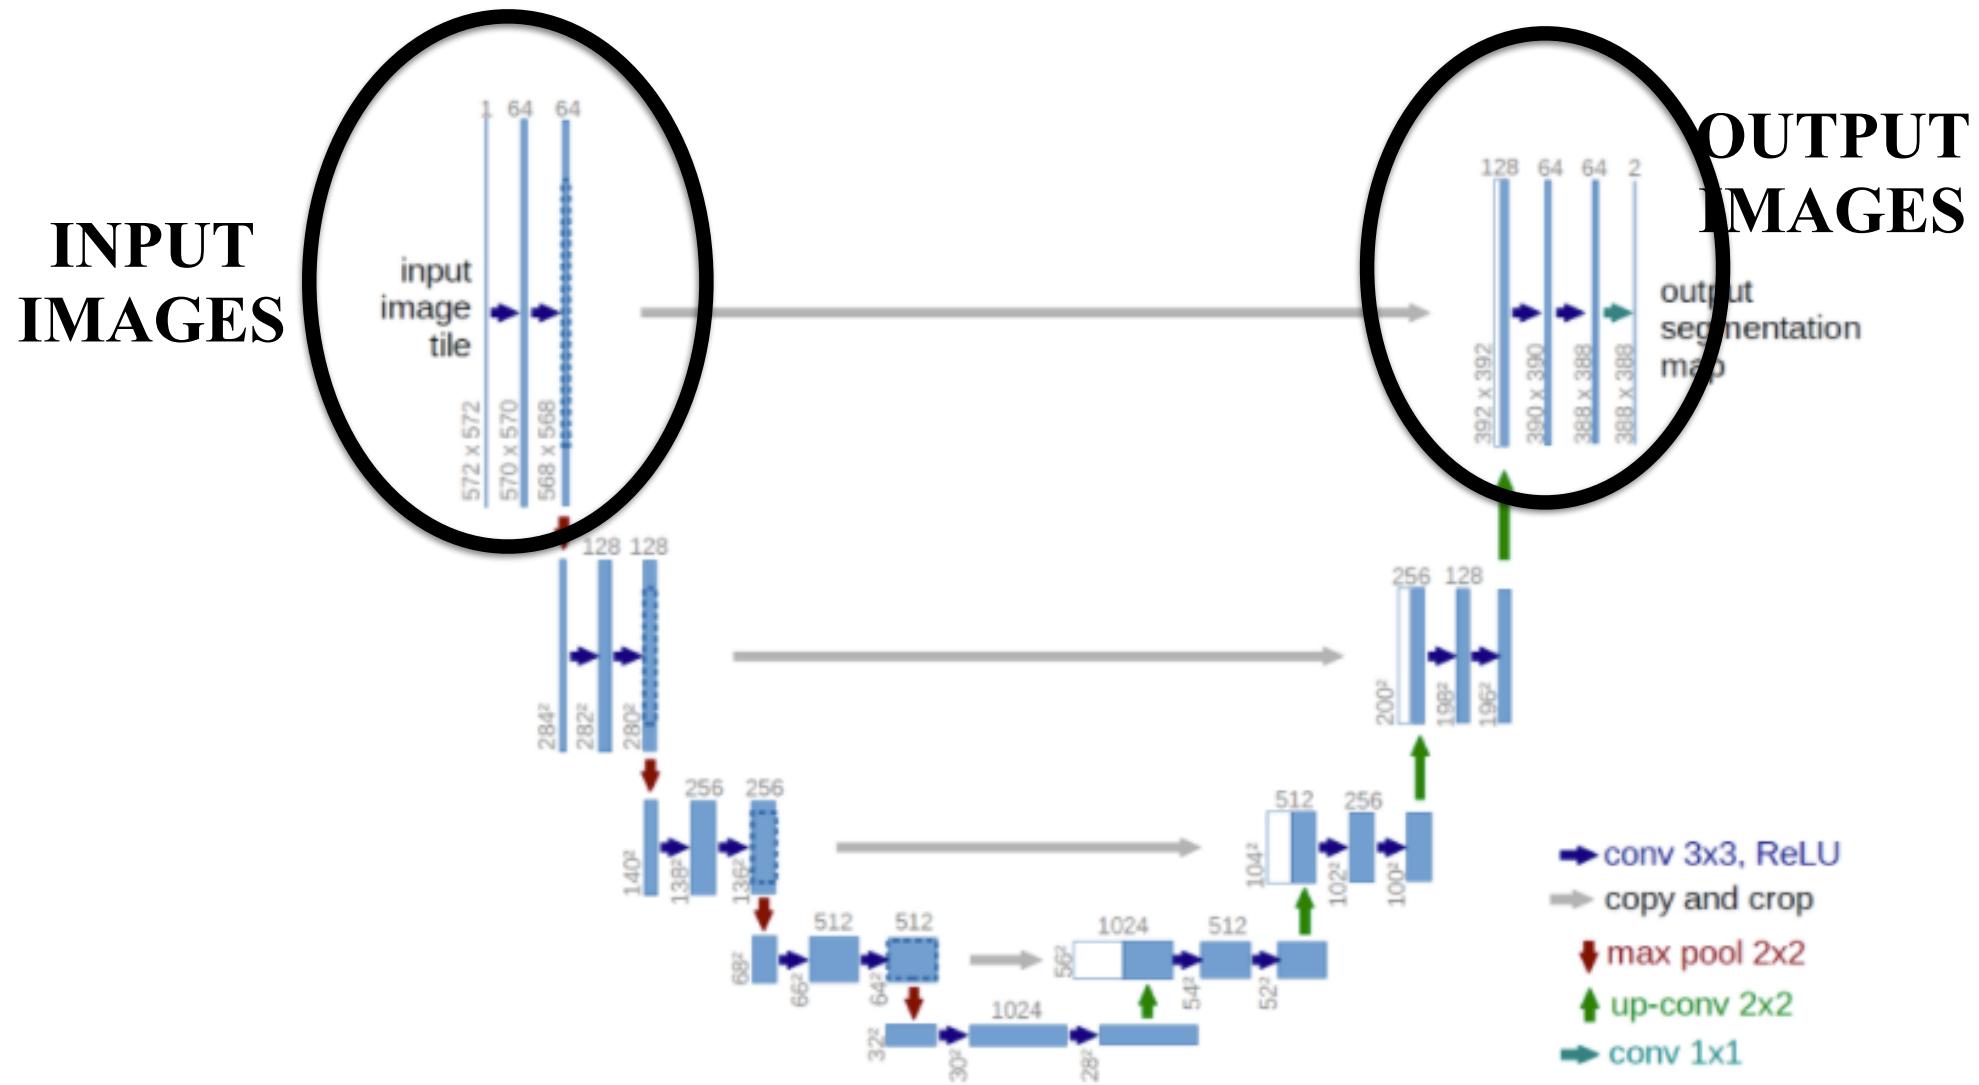
DECODER-ENCODER NETWORKS



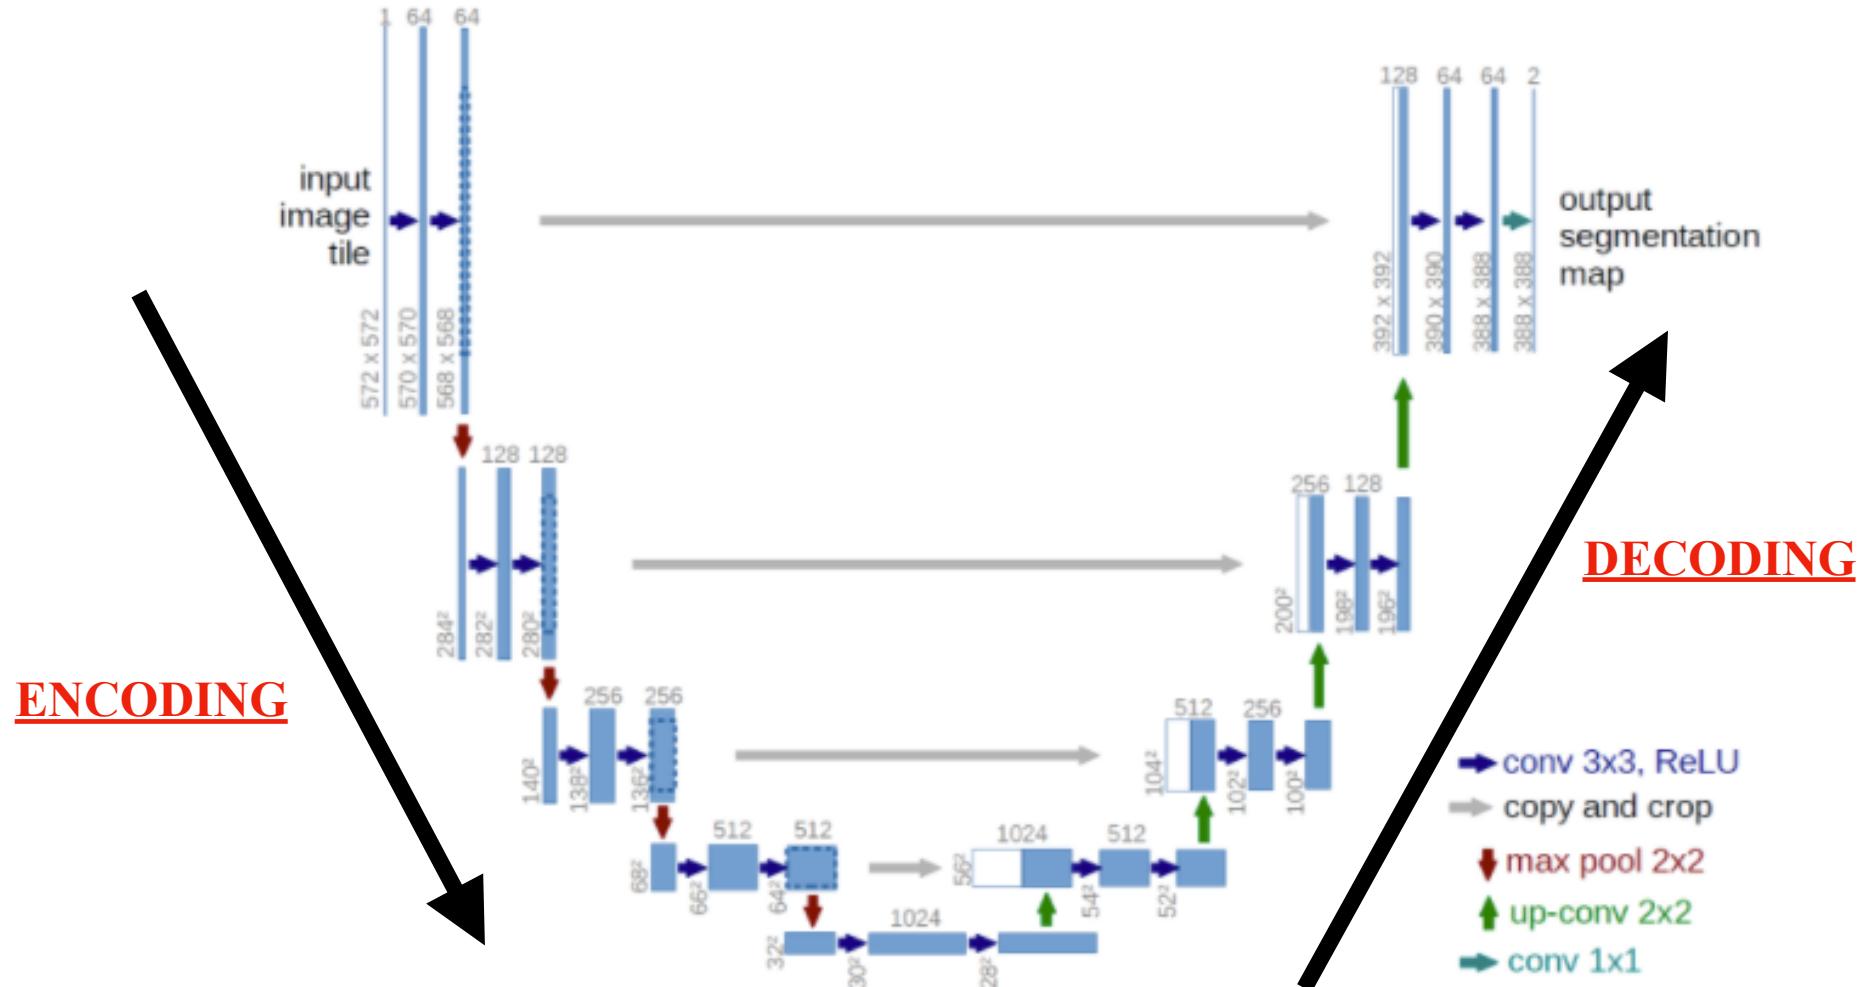
MAIN APPLICATION IS IMAGE SEGMENTATION



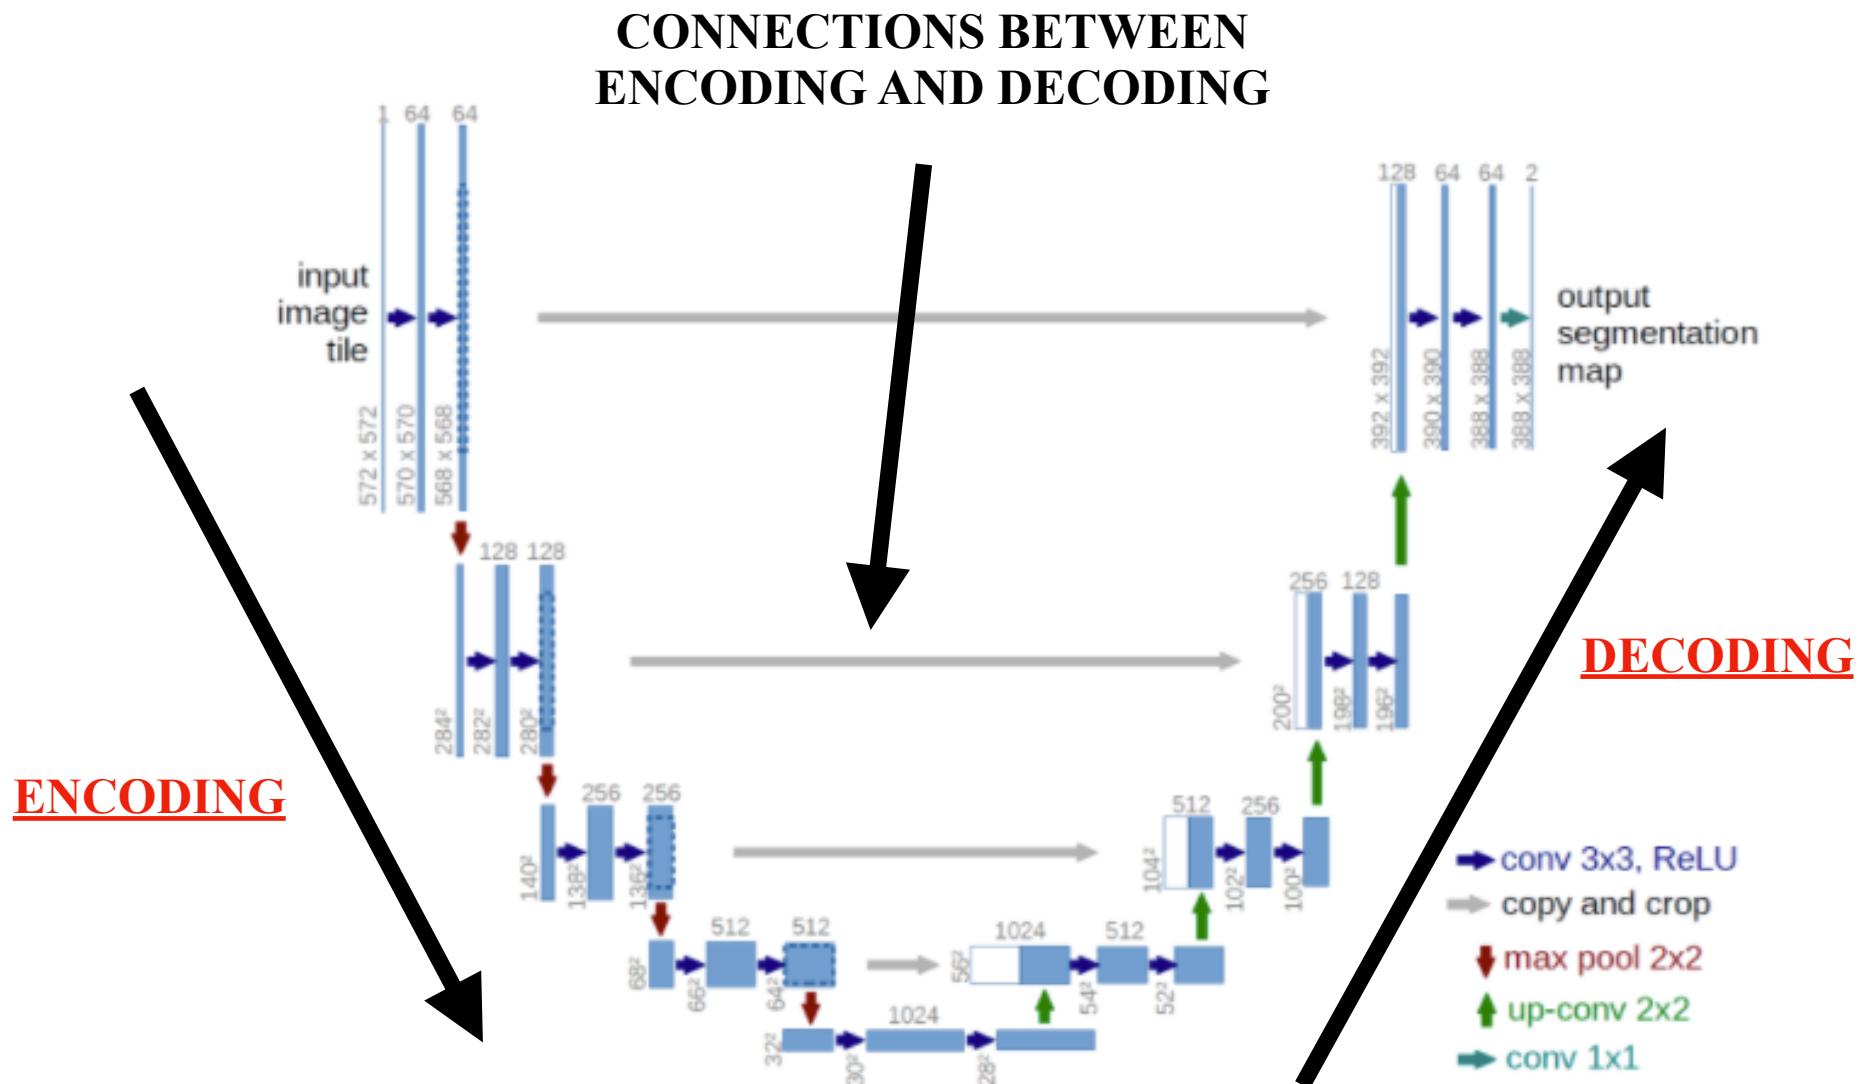
ENCODING-DECODING TO EXTRACT IMAGE FEATURES: U-NET



ENCODING-DECODING TO EXTRACT IMAGE FEATURES: U-NET



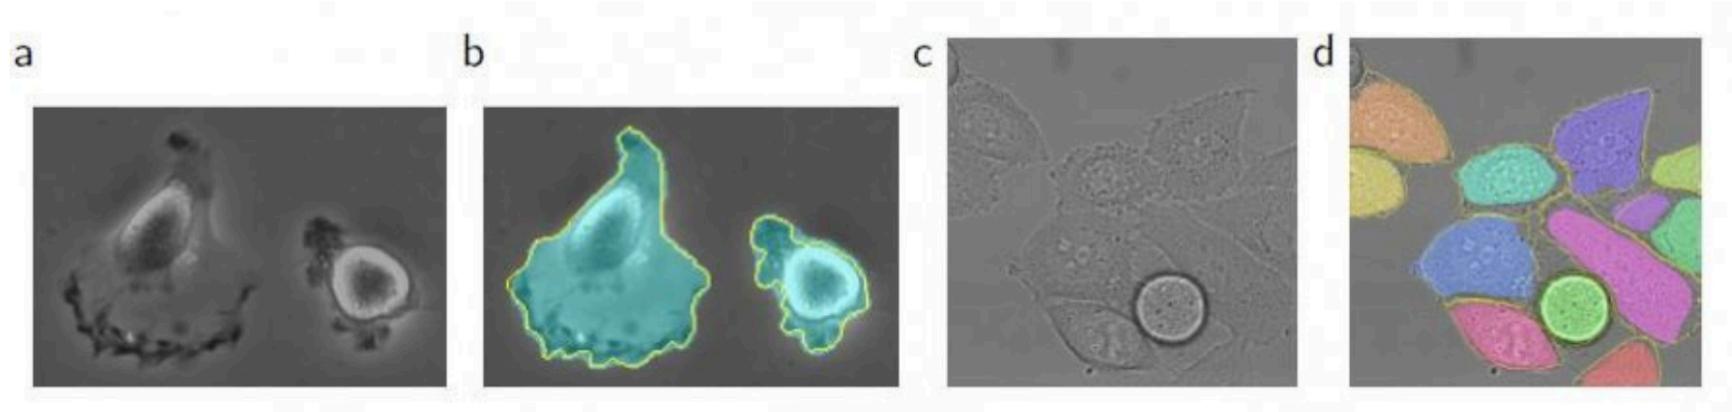
ENCODING-DECODING TO EXTRACT IMAGE FEATURES: U-NET



WHAT IS THIS USEFUL
FOR?

WHAT IS THIS USEFUL FOR?

A STANDARD TOOL FOR IMAGE SEGMENTATION

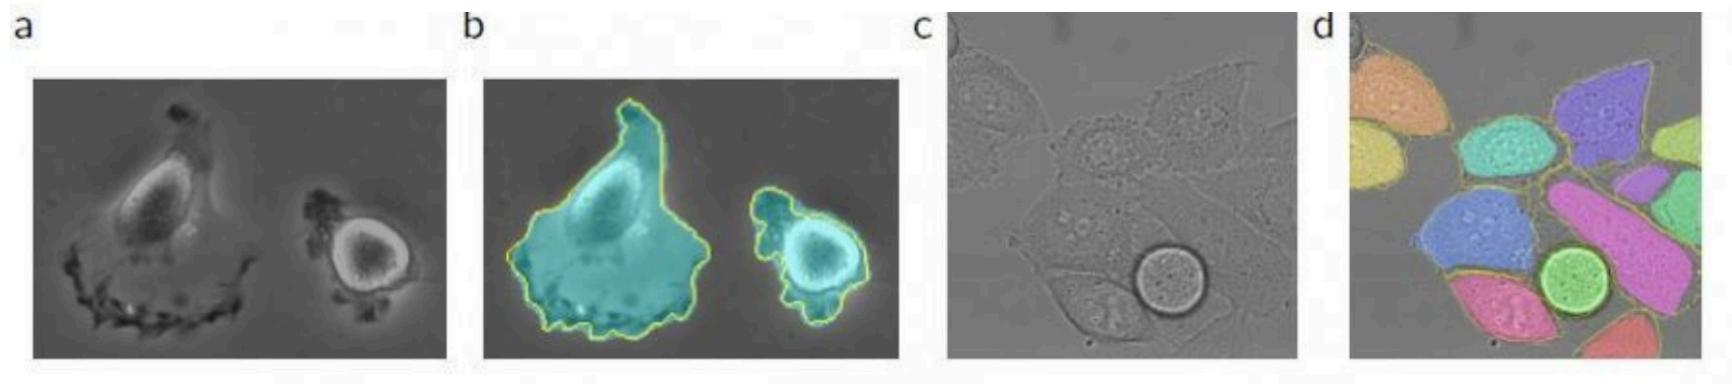


Result on the ISBI cell tracking challenge. (a) part of an input image of the PhC-U373 data set. (b) Segmentation result (cyan mask) with the manual ground truth (yellow border) (c) input image of the DIC-HeLa data set. (d) Segmentation result (random colored masks) with the manual ground truth (yellow border).

WHAT IS THIS USEFUL FOR?

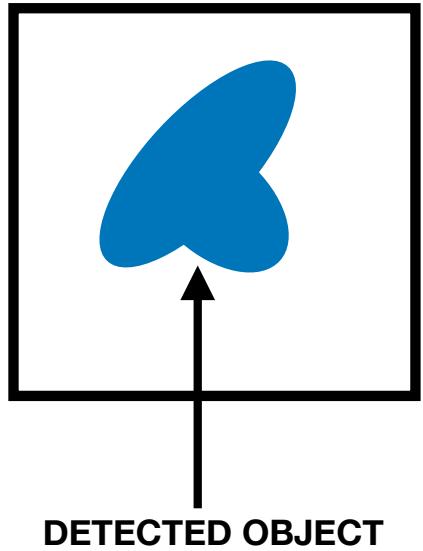
A STANDARD TOOL FOR IMAGE SEGMENTATION

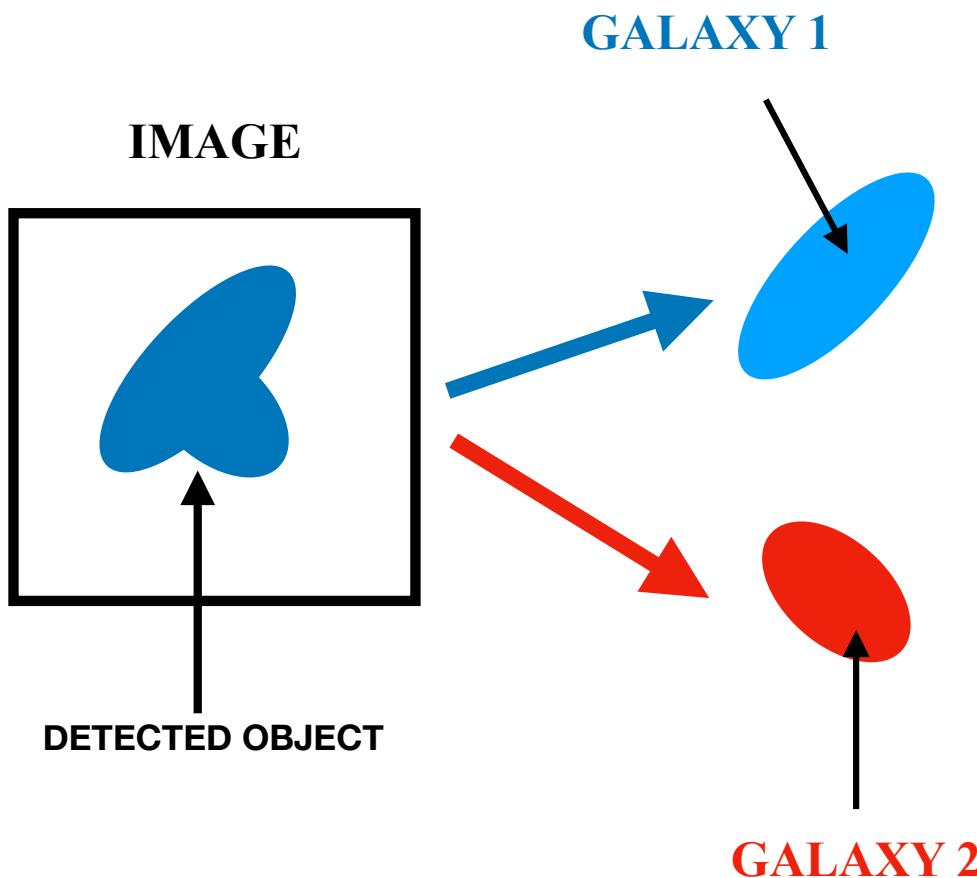
MORE ON THE SEMINAR LATER TODAY

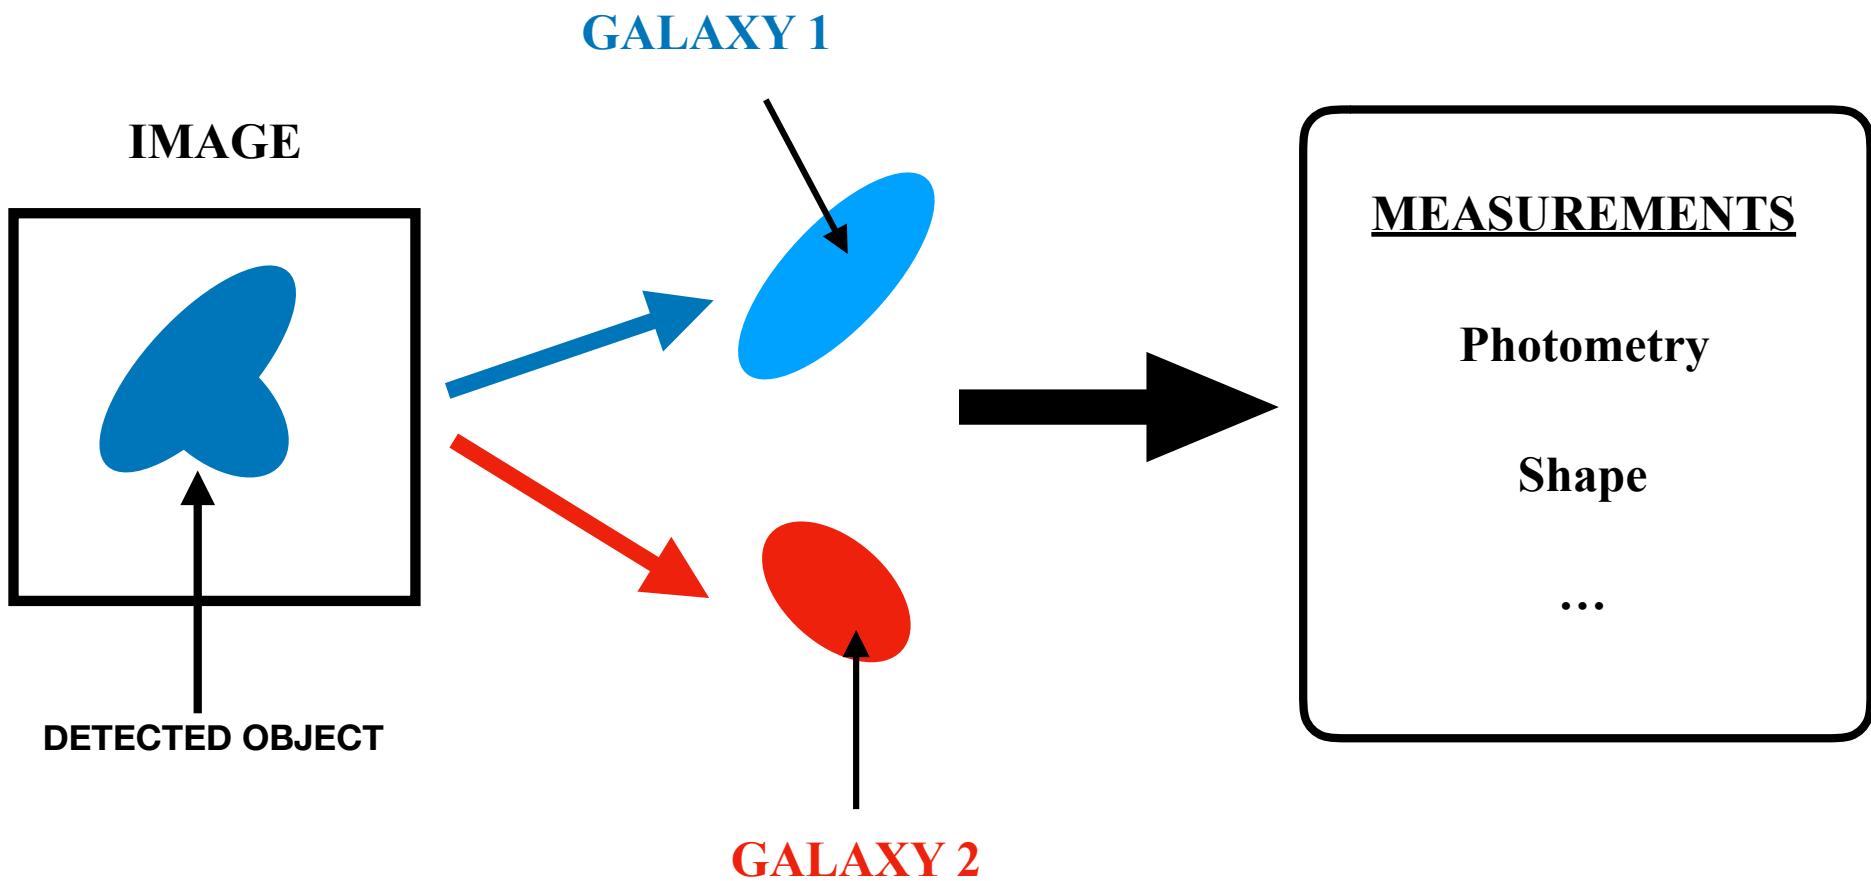


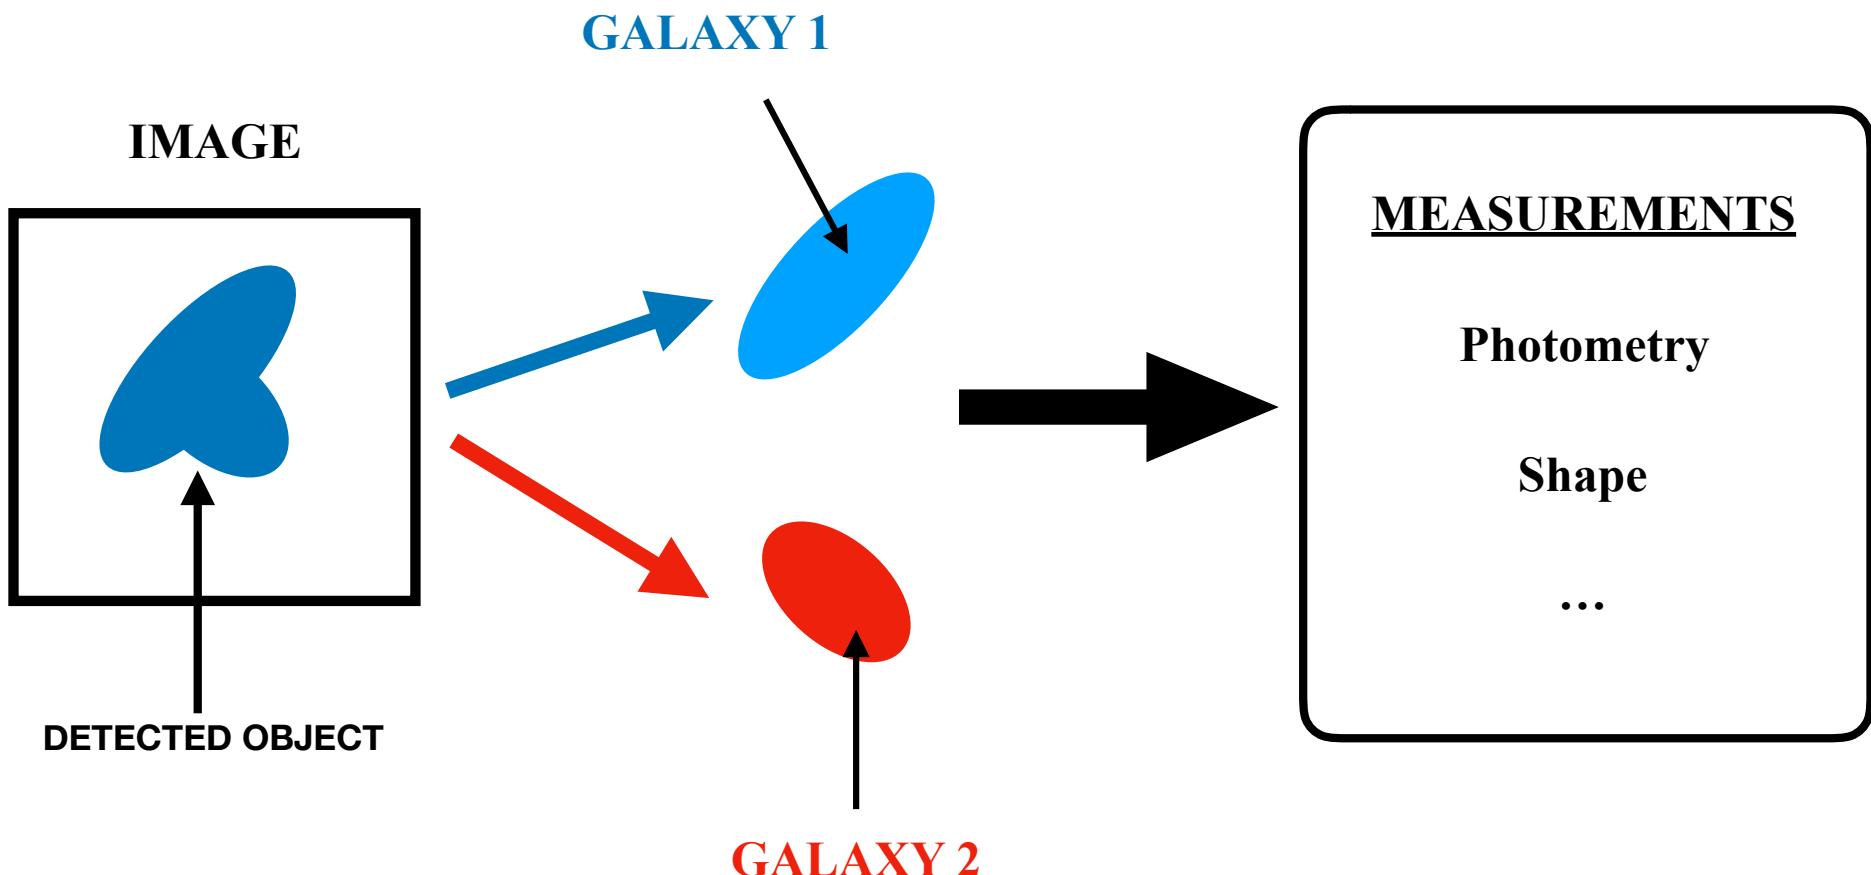
Result on the ISBI cell tracking challenge. (a) part of an input image of the PhC-U373 data set. (b) Segmentation result (cyan mask) with the manual ground truth (yellow border) (c) input image of the DIC-HeLa data set. (d) Segmentation result (random colored masks) with the manual ground truth (yellow border).

IMAGE

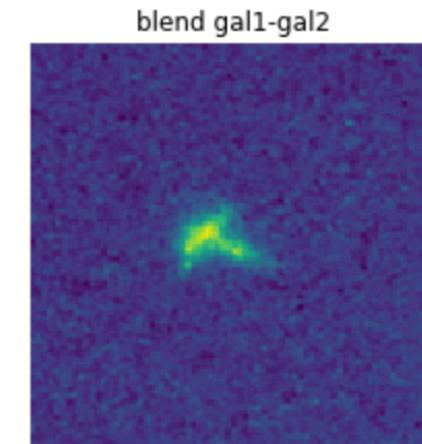
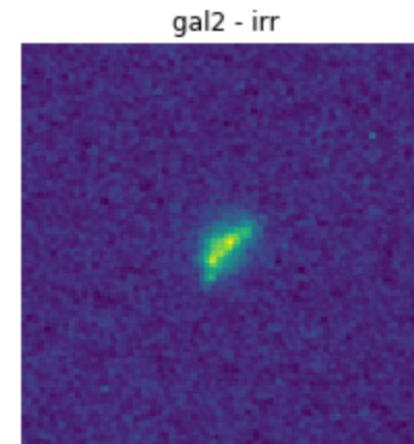
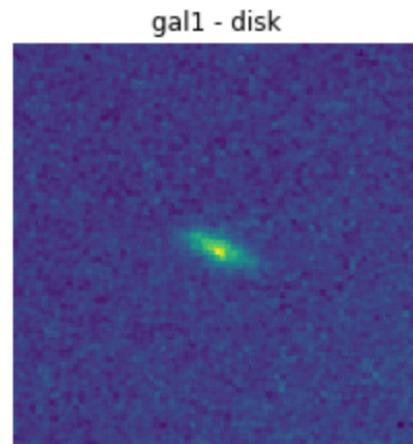


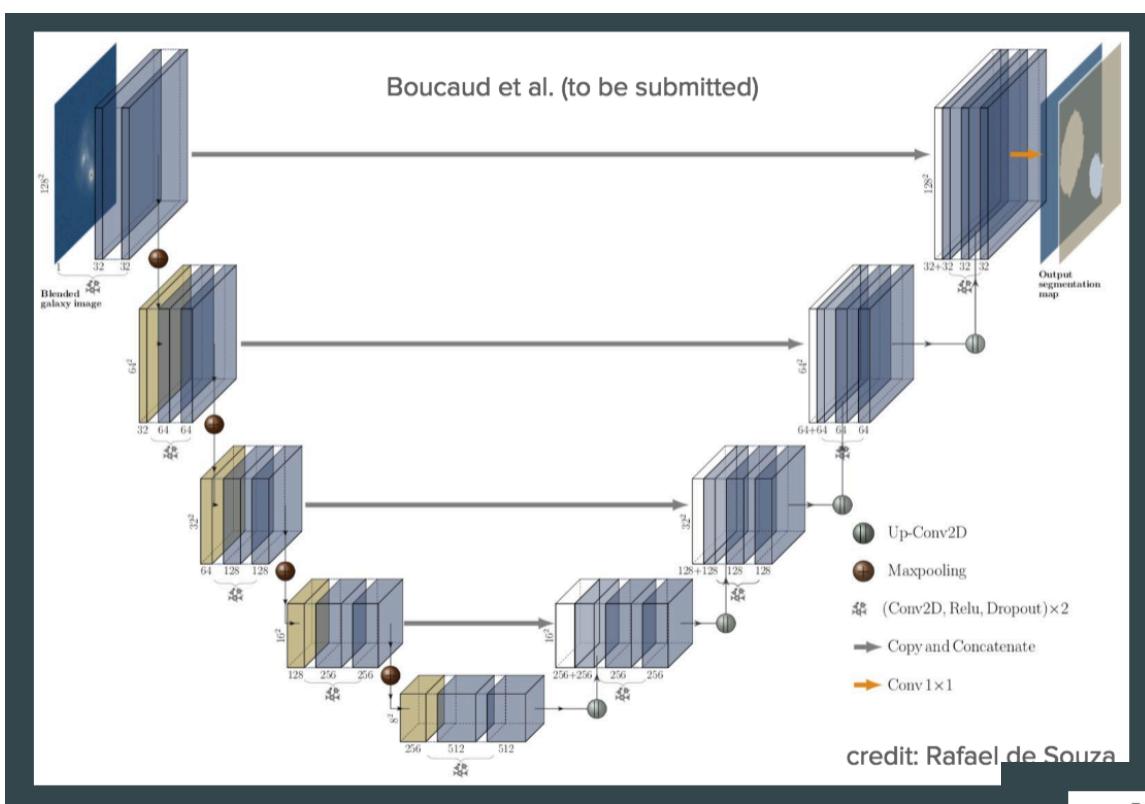




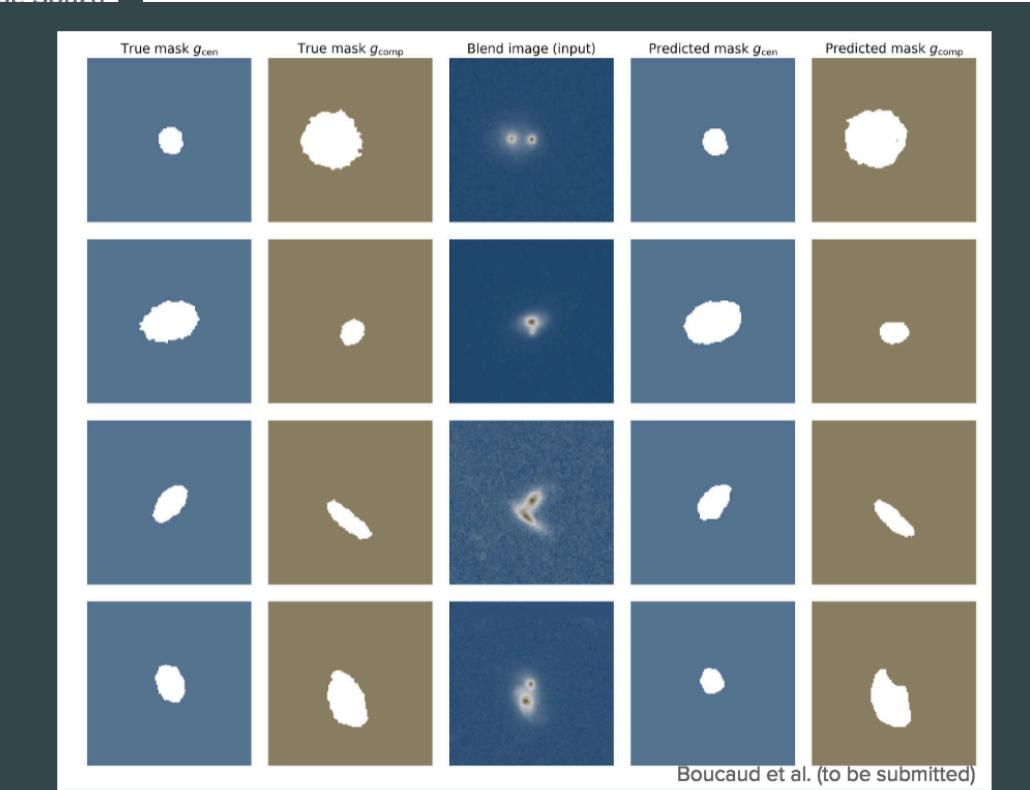


**ISOLATED
GALAXIES
ARTIFICIALLY
BLENDED**



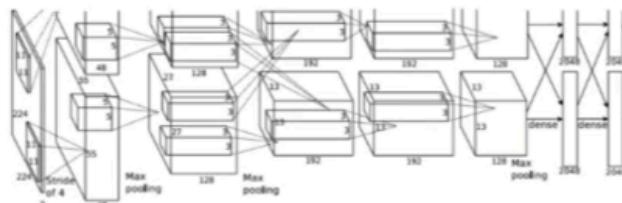


credit: Rafael de Souza

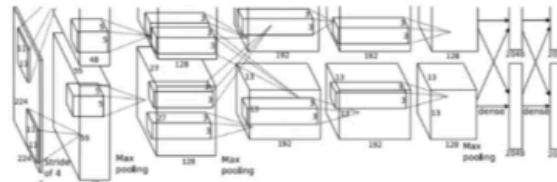


A WORD ON R-CNNs..

JUST FOR YOUR RECORDS...



CAT: (x, y, w, h)



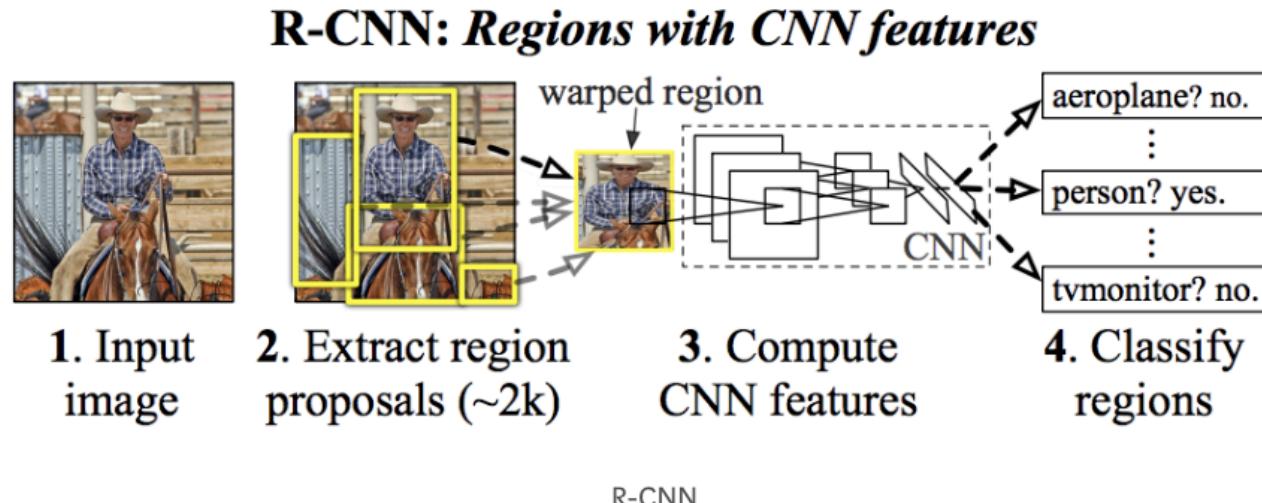
DUCK: (x, y, w, h)
DUCK: (x, y, w, h)

....

Girshick+14

A WORD ON R-CNNs..

JUST FOR YOUR RECORDS...

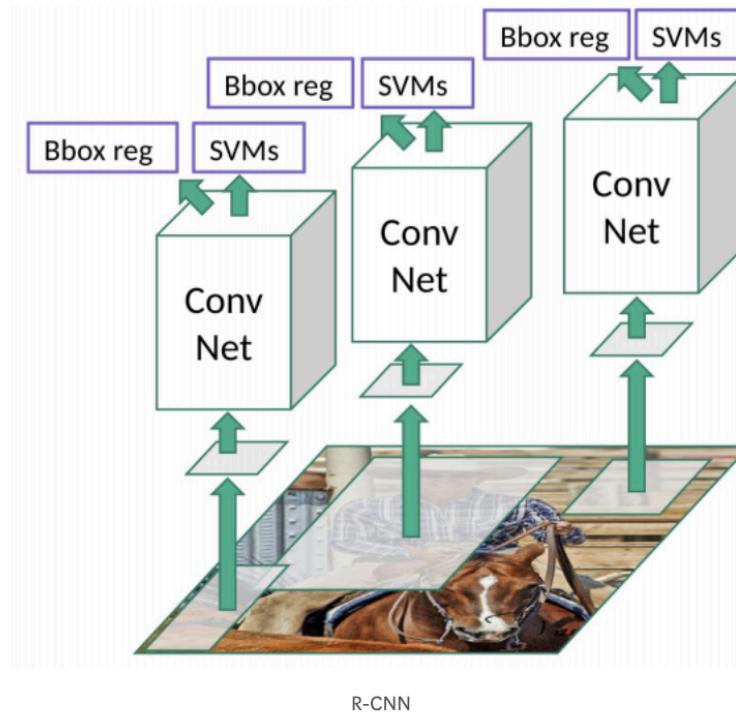


A WORKD ON R-CNNs..

JUST FOR YOUR RECORDS...

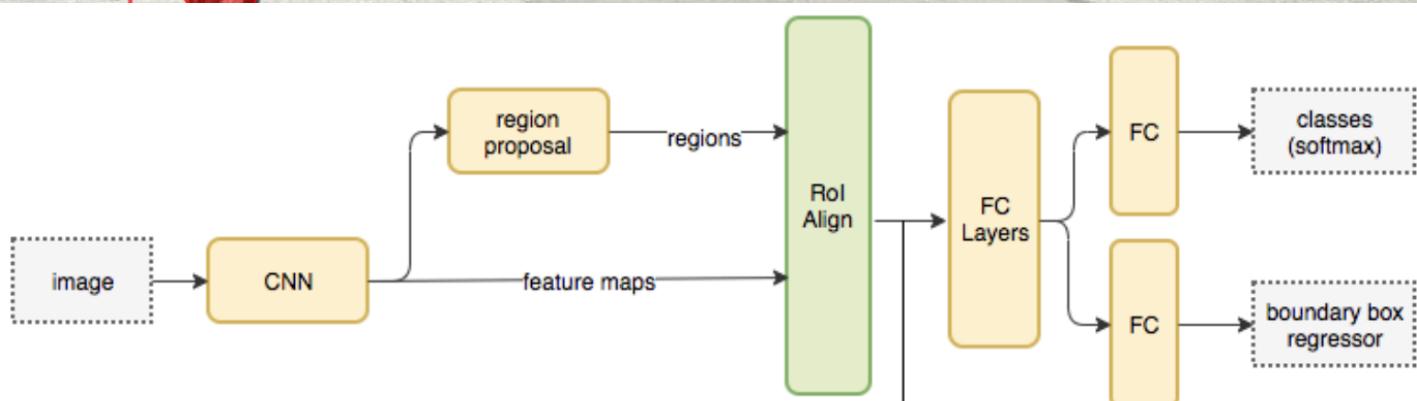
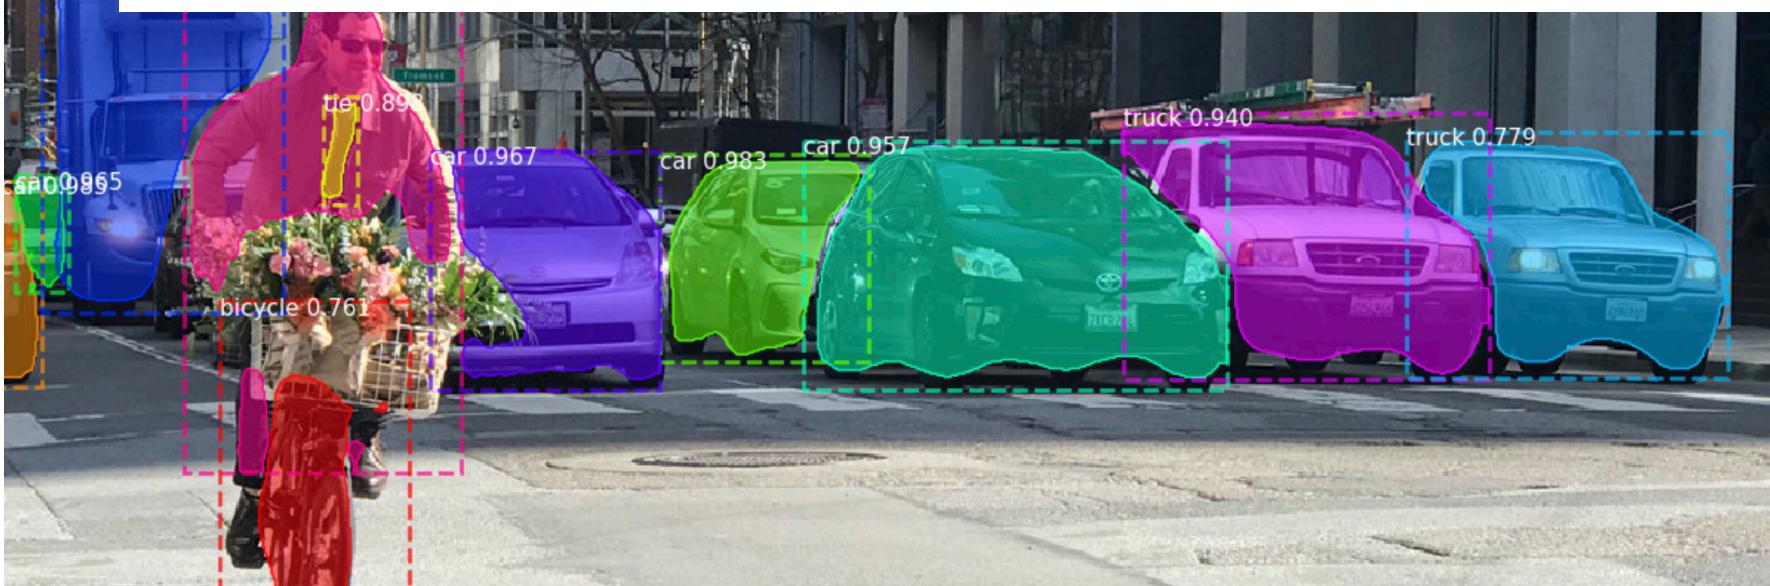


1. Input image



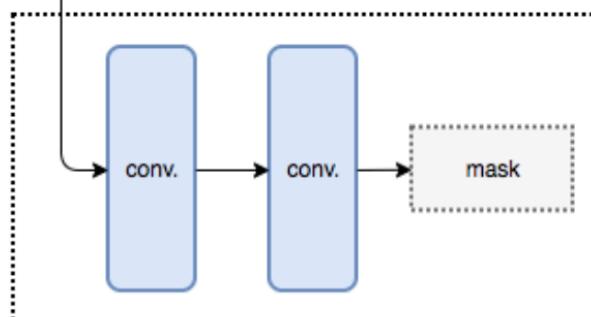
eroplane? no.
⋮
erson? yes.
⋮
vmonitor? no.
. Classify regions

SIMULTANEOUS DETECTION + SEGMENTATION + CLASSIFICATION



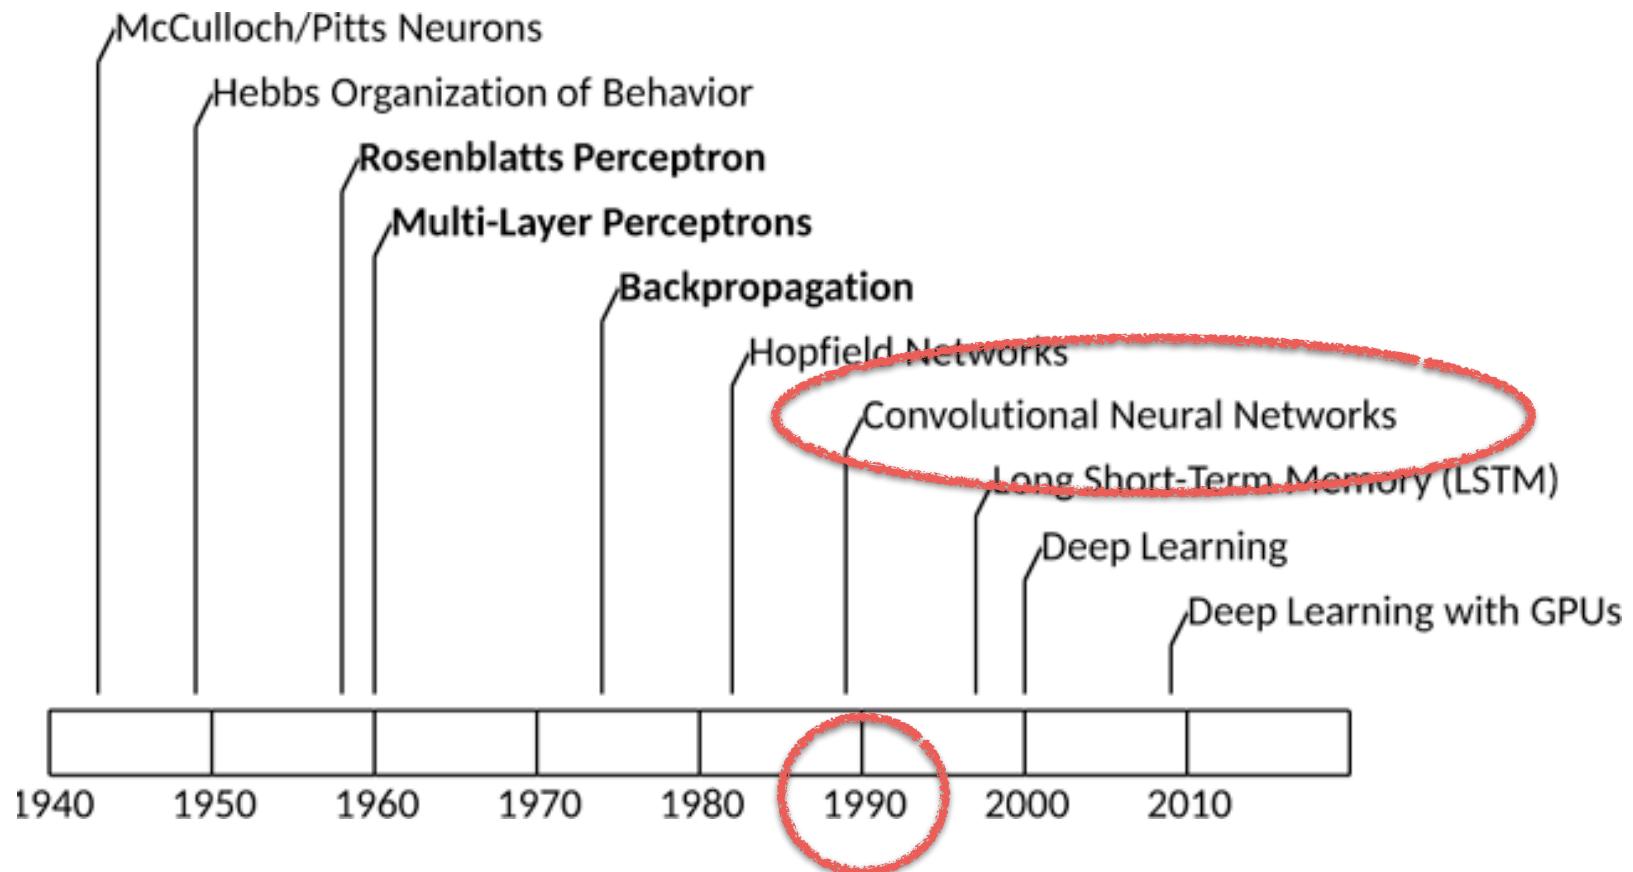
MASK R-CNN

He+17



Mask

WELL, BUT THIS IS AN “OLD” IDEA - WHY NOW?



WELL, BUT THIS IS AN “OLD” IDEA - WHY NOW?

1 - MORE DATA TO TRAIN! DEEP NETWORKS HAVE A
LARGE NUMBER OF PARAMETERS - THX TO SOCIAL
MEDIA ...

WELL, BUT THIS IS AN
“OLD” IDEA - WHY NOW?

2 - GPUs - TRAINING OF THESE DEEP NETWORKS
HAS REMAINED PROHIBITIVELY TIME CONSUMING
WITH CPUs - THX TO VIDEO GAMES...

GPUs vs. CPUs

CPUs

FEWER CORES (~10x)

EACH CORE IS FASTER

USEFUL FOR
SEQUENTIAL TASKS

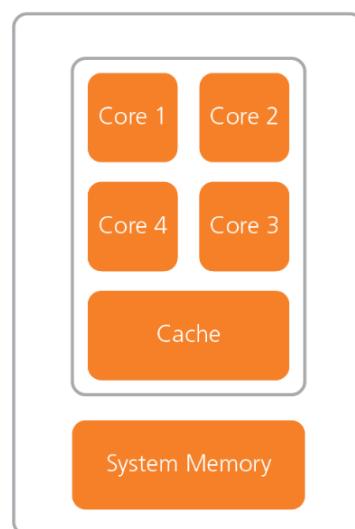
GPUs

MORE CORES (100x)

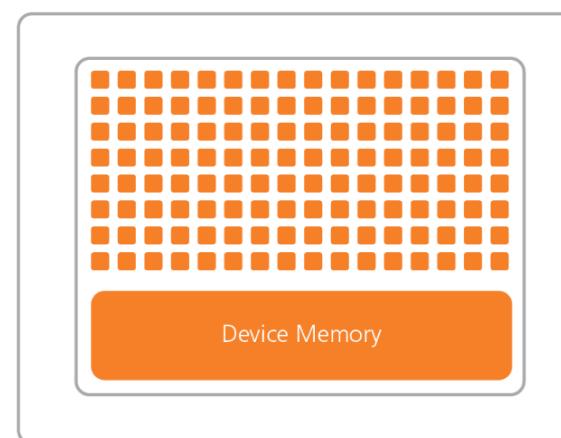
EACH CORE IS SLOWER

USEFUL FOR PARALLEL
TASKS

CPU (Multiple Cores)



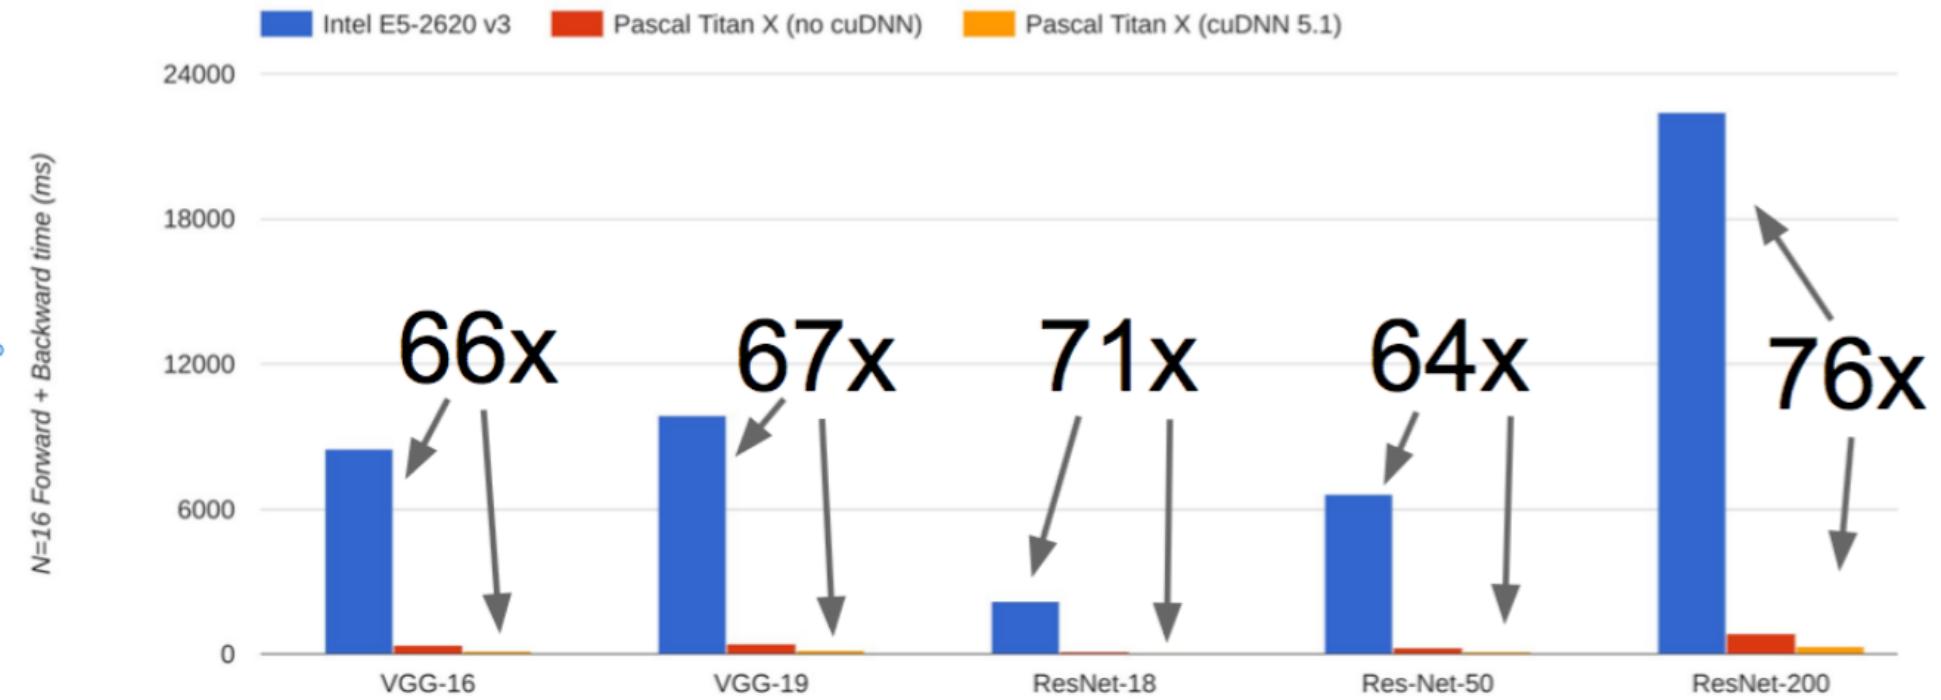
GPU (Hundreds of Cores)



Slide Credit:

GPUs vs. CPUs

More benchmarks available [here](#).



GPUs for deep learning

NVIDIA GPUs ARE PROGRAMMED THROUGH CUDA
[Compute Unified Device Architecture]

ANOTHER ALTERNATIVE IS OPENCL, SUPPORTED BY
SEVERAL MANUFACTURES, LESS INVESTMENT [Way less
used]

CuDNN IS A LIBRARY FOR SPECIFIC DEEP LEARNING
COMPUTATIONS ON NVIDIA GPUs

THE PRICE TO PAY?

1. LARGE NUMBER OF PARAMETERS IMPLIES LARGE DATASETS TO TRAIN
2. LOOSE EVEN MORE DEGREE OF CONTROL OF WHAT THE ALGORITHM IS DOING SINCE THE FEATURE EXTRACTION PROCESS BECOMES UNSUPERVISED

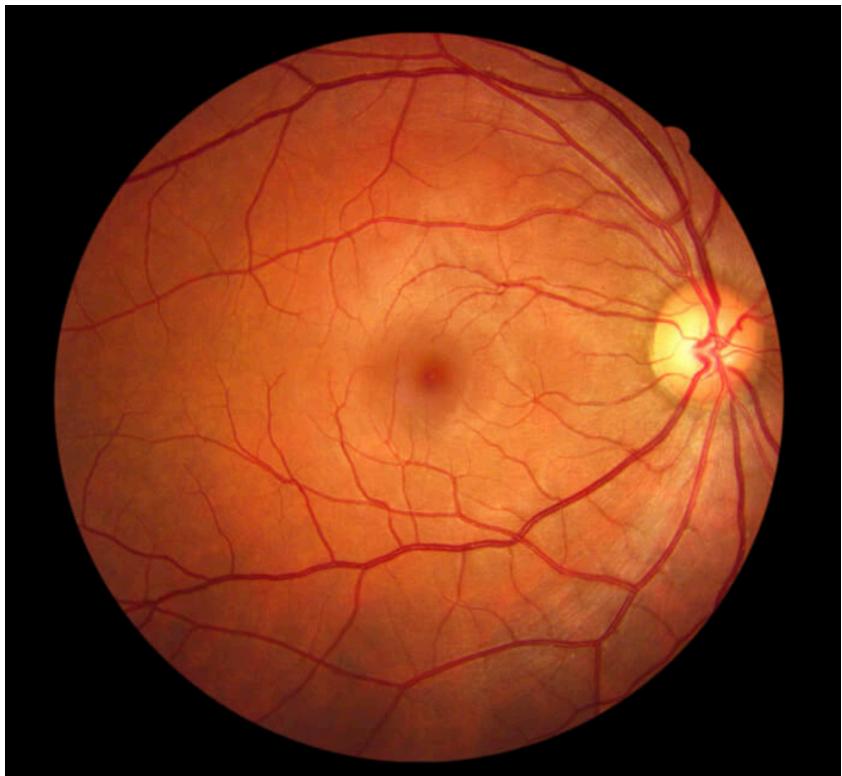
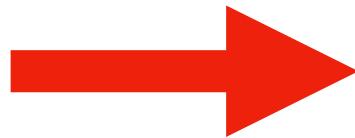
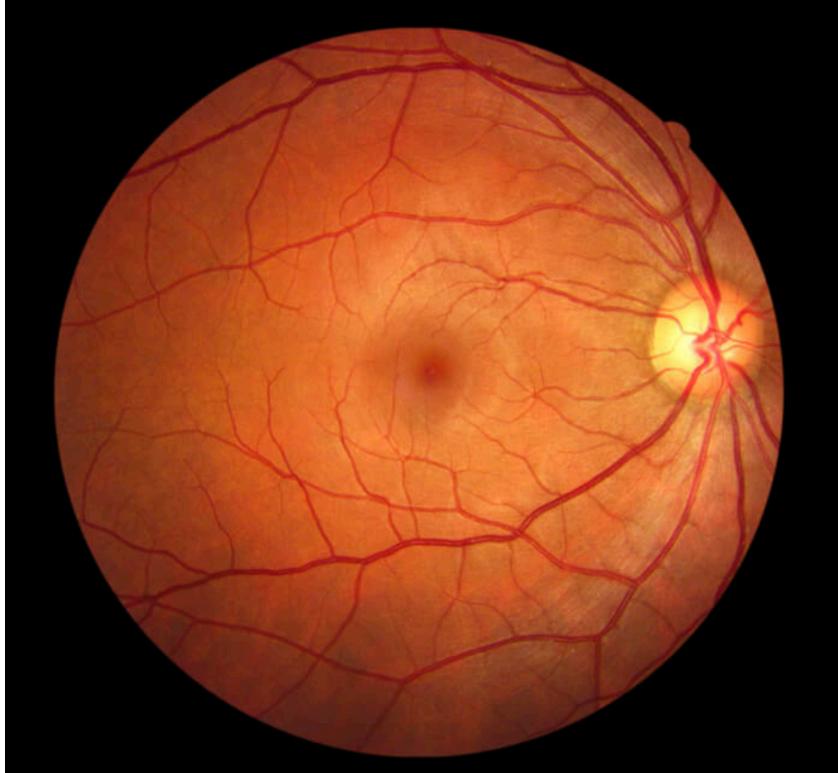


IMAGE OF THE BACK OF THE EYE





**DEEP LEARNING CAN
IDENTIFY
THE PATIENT'S
GENDER WITH 95%
ACCURACY**

IMAGE OF THE BACK OF THE EYE



VISUALIZING CNNs

[understanding CNNs decisions]

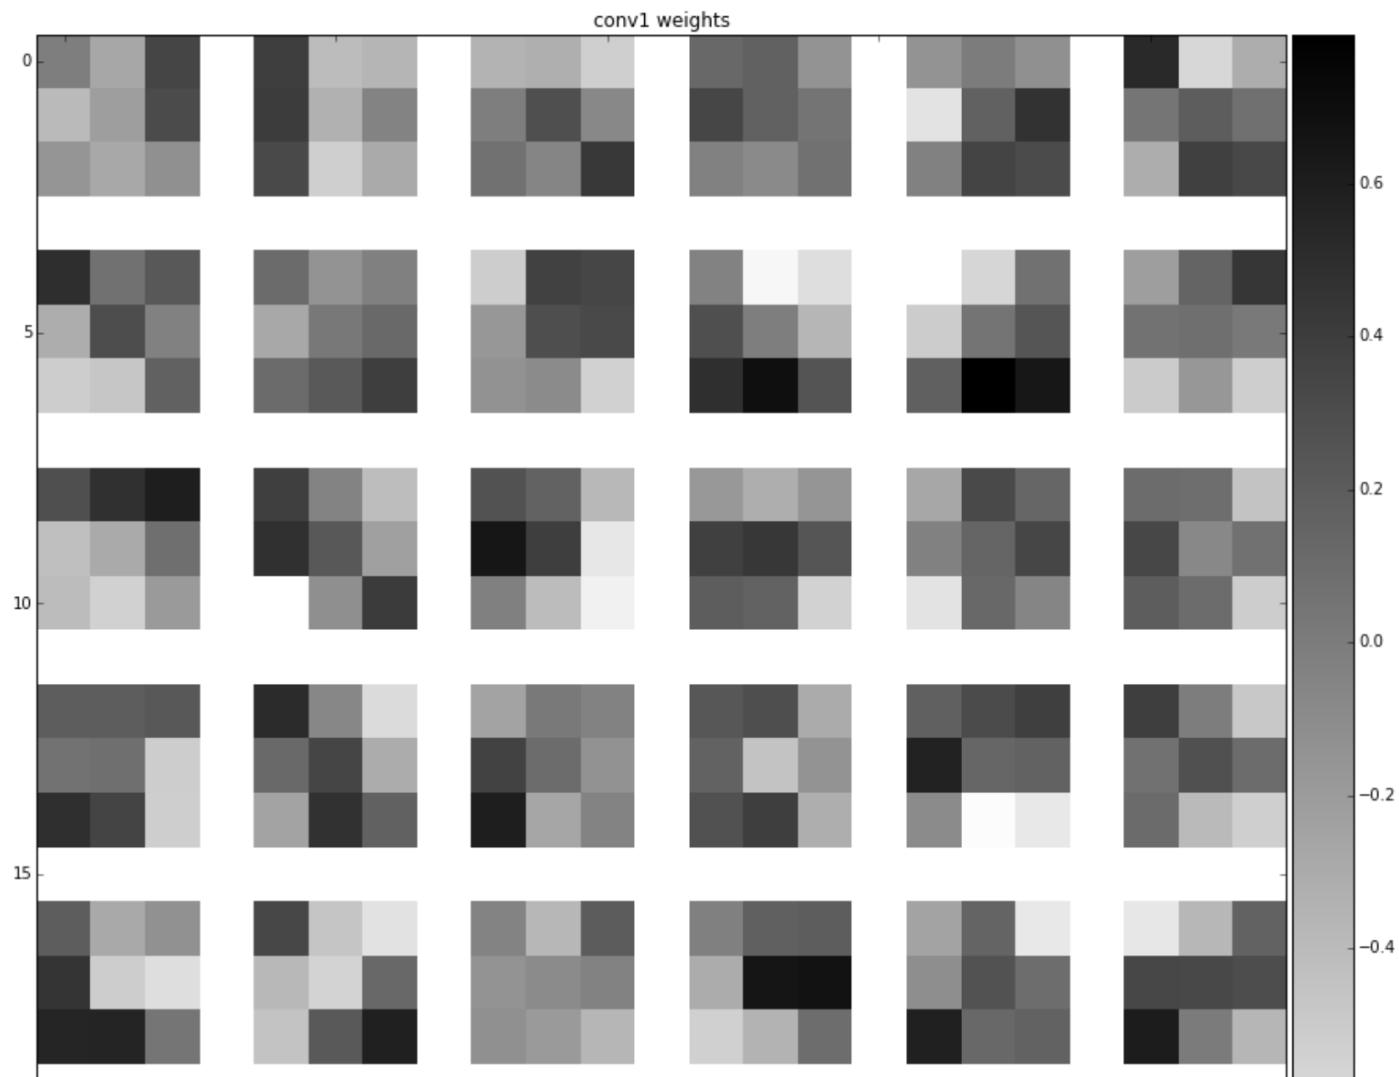
Attribution techniques

DEEP NETWORKS ARE “BLACK BOXES”?

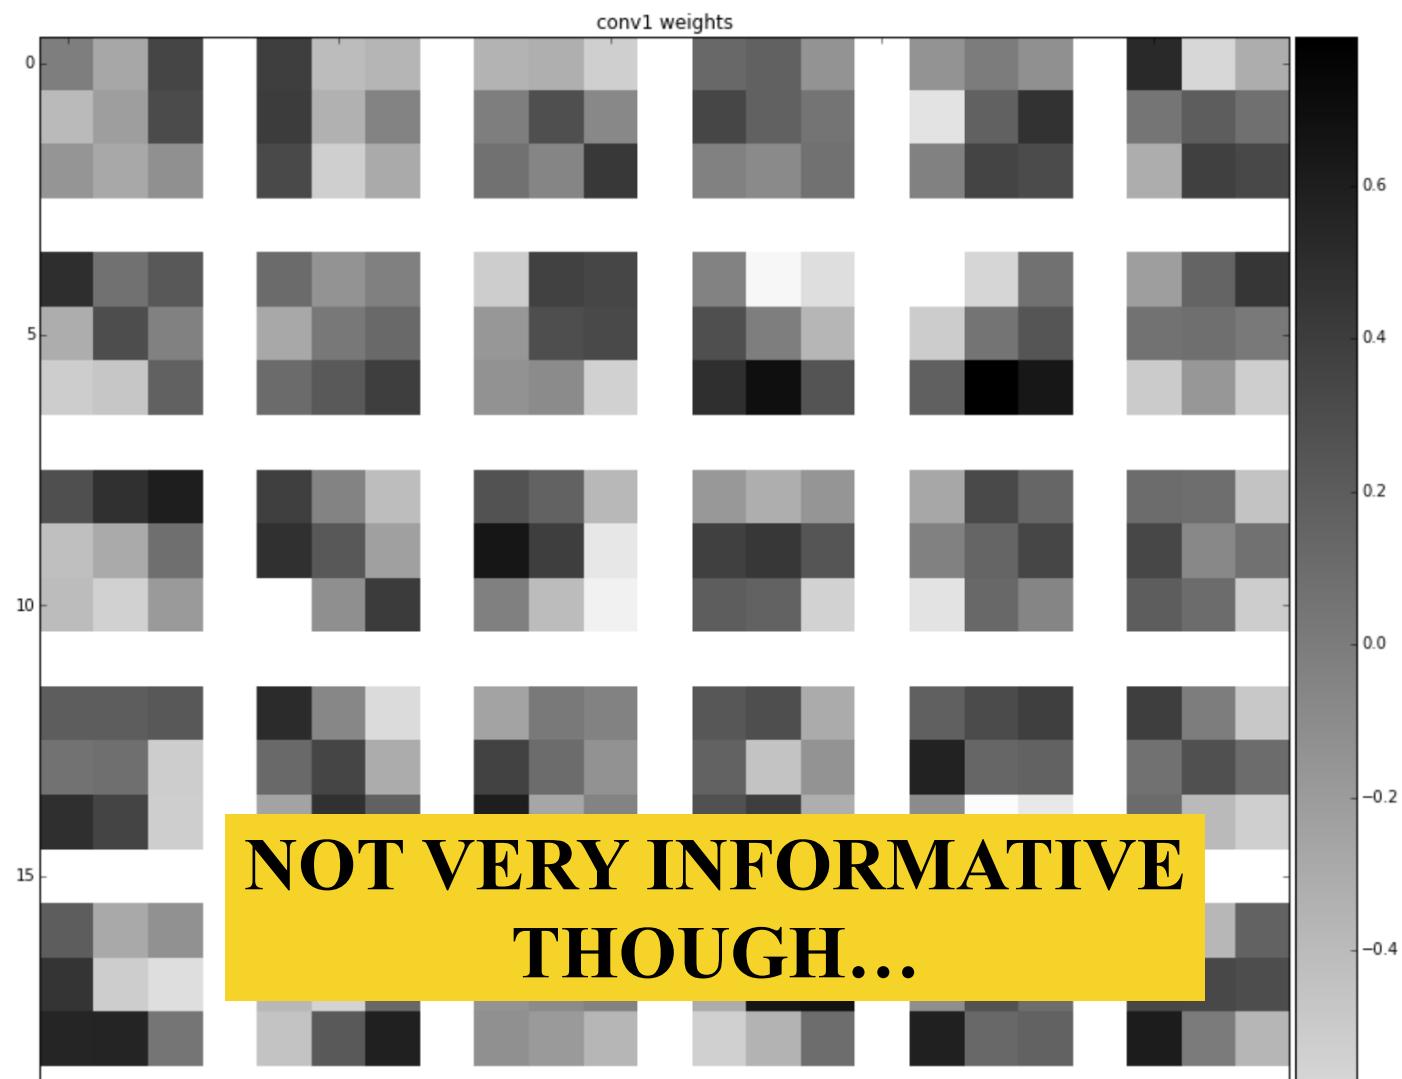
INTERPRETING THE RESULTS IS
EXTREMELY DIFFICULT

THIS IS TRUE BUT A LOT OF WORK
IS DONE TO UNVEIL THEIR BEHAVIOR

THE SIMPLEST APPROACH IS TO VISUALIZE THE LEARNED WEIGHTS AT INTERMEDIATE LAYERS



THE SIMPLEST APPROACH IS TO VISUALIZE THE LEARNED WEIGHTS AT INTERMEDIATE LAYERS



THE SIMPLEST APPROACH IS TO VISUALIZE THE LEARNED WEIGHTS AT INTERMEDIATE LAYERS

IN KERAS:

```
# build model
model = Sequential()
model.add(Convolution2D(depth, conv_size0, conv_size0, activation=act,
border_mode='same', name = "conv0",
                     input_shape=(img_channels, img_rows, img_cols),
                     init=initialization, W_constraint=constraint))
model.add(Dropout(dropout_rate_conv))

# get the symbolic outputs of each "key" layer (we gave them unique names).
layer_dict = dict([(layer.name, layer) for layer in model.layers])

layer_dict[layer_name].W.get_value(borrow=True)
W = np.squeeze(W)
print("W shape : ", W.shape)

# plot weights
pl.figure(figsize=(15, 15))
pl.title('conv1 weights')
nice_imshow(pl.gca(), make_mosaic(W, 6, 6), cmap=cm.binary)
```

THE SIMPLEST APPROACH IS TO VISUALIZE THE LEARNED WEIGHTS AT INTERMEDIATE LAYERS

IN KERAS:

give names to layers

```
# build model
model = Sequential()
model.add(Convolution2D(depth, conv_size0, conv_size0, activation=act,
border_mode='same', name = "conv0",
input_shape=(img_channels, img_rows, img_cols),
init=initialization, W_constraint=constraint))
model.add(Dropout(dropout_rate_conv))

# get the symbolic outputs of each "key" layer (we gave them unique names).
layer_dict = dict([(layer.name, layer) for layer in model.layers])

layer_dict[layer_name].W.get_value(borrow=True)
W = np.squeeze(W)
print("W shape : ", W.shape)

# plot weights
pl.figure(figsize=(15, 15))
pl.title('conv1 weights')
nice_imshow(pl.gca(), make_mosaic(W, 6, 6), cmap=cm.binary)
```

THE SIMPLEST APPROACH IS TO VISUALIZE THE LEARNED WEIGHTS AT INTERMEDIATE LAYERS

IN KERAS: create dictionary to link layers to names

```
# build model
model = Sequential()
model.add(Convolution2D(depth, conv_size0, conv_size0, activation=act,
border_mode='same', name = "conv0",
                     input_shape=(img_channels, img_rows, img_cols),
                     init=initialization, W_constraint=constraint))
model.add(Dropout(dropout_rate_conv))

# get the symbolic outputs of each "key" layer (we gave them unique names).
layer_dict = dict([(layer.name, layer) for layer in model.layers])

layer_dict[layer_name].W.get_value(borrow=True)
W = np.squeeze(W)
print("W shape : ", W.shape)

# plot weights
pl.figure(figsize=(15, 15))
pl.title('conv1 weights')
nice_imshow(pl.gca(), make_mosaic(W, 6, 6), cmap=cm.binary)
```

THE SIMPLEST APPROACH IS TO VISUALIZE THE LEARNED WEIGHTS AT INTERMEDIATE LAYERS

IN KERAS:

for a given name, get the weights

```
# build model
model = Sequential()
model.add(Convolution2D(depth, conv_size0, conv_size0, activation=act,
border_mode='same', name = "conv0",
                     input_shape=(img_channels, img_rows, img_cols),
                     init=initialization, W_constraint=constraint))
model.add(Dropout(dropout_rate_conv))

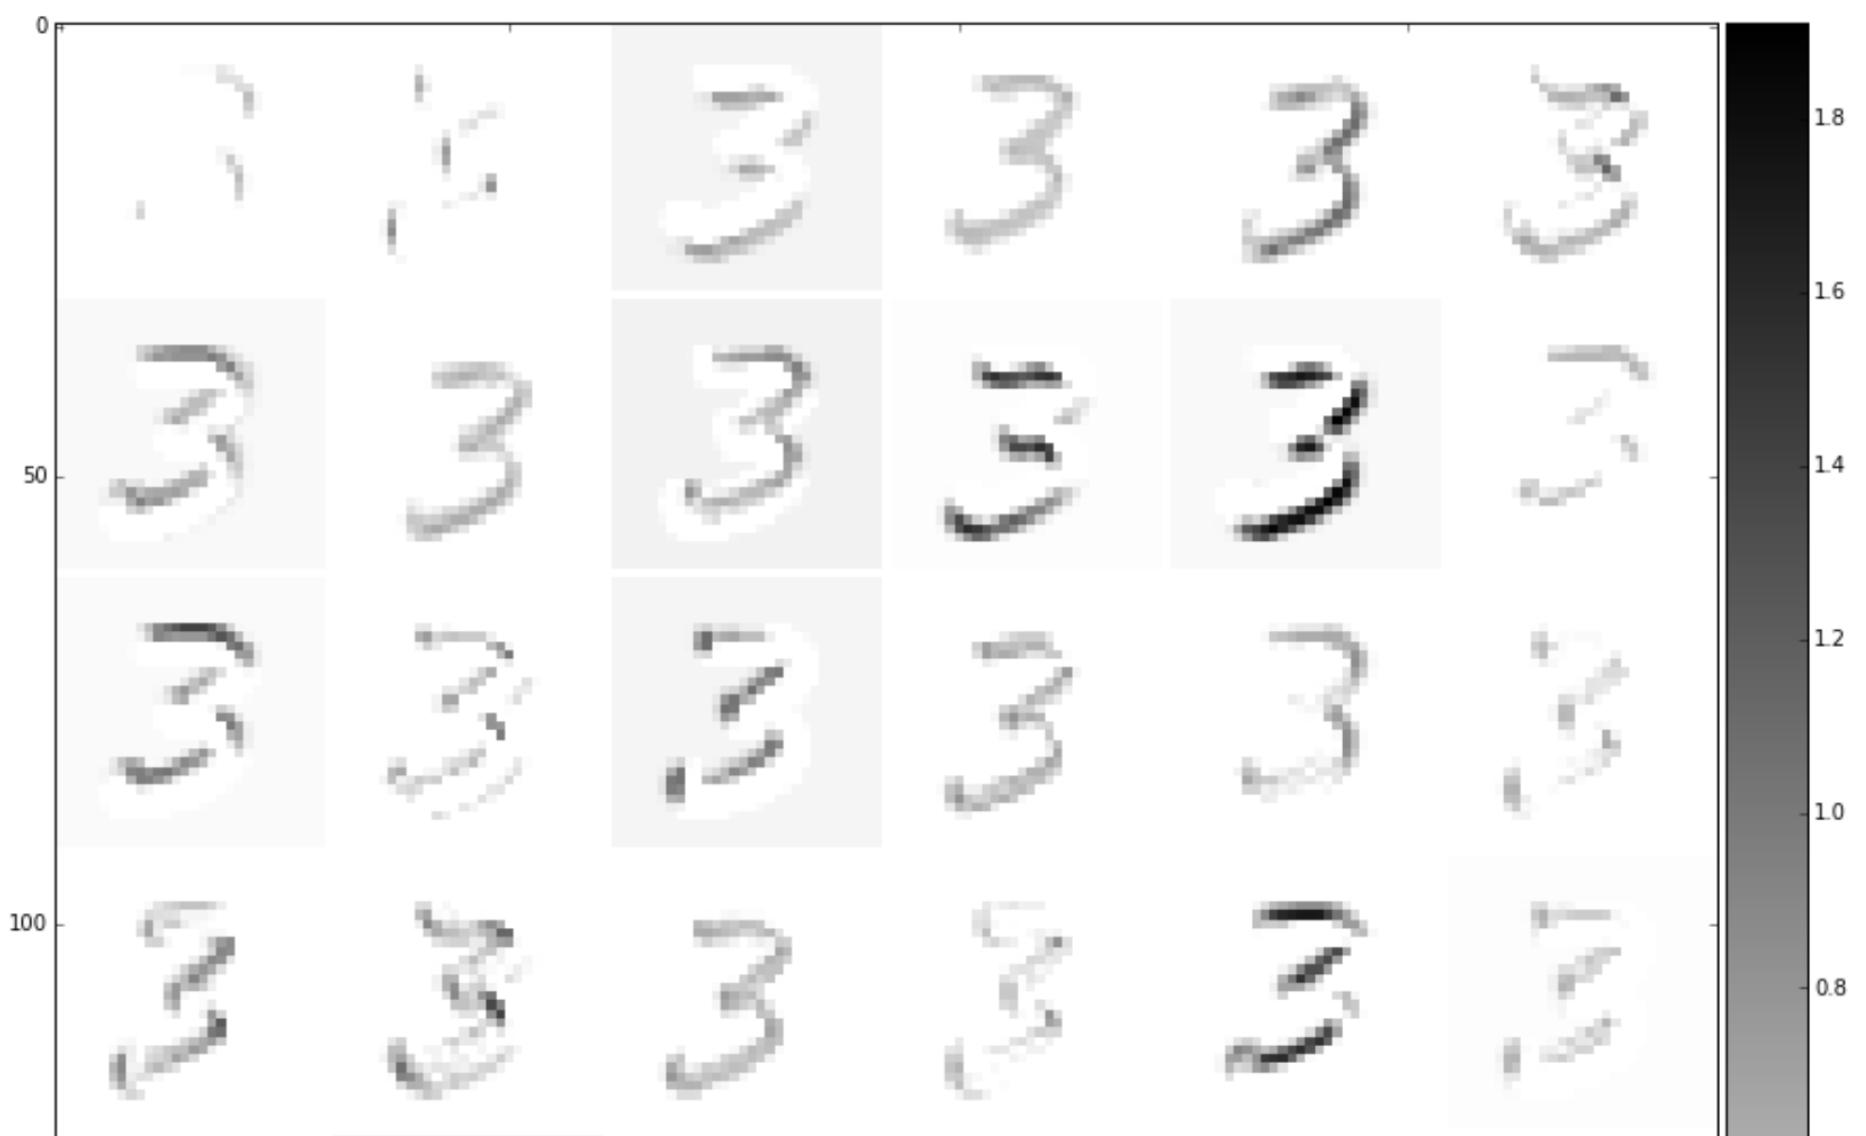
# get the symbolic outputs of each "key" layer (we gave them unique names).
layer_dict = dict([(layer.name, layer) for layer in model.layers])

layer_dict[layer_name].W.get_value(borrow=True)
W = np.squeeze(W)
print("W shape : ", W.shape)

# plot weights
pl.figure(figsize=(15, 15))
pl.title('conv1 weights')
nice_imshow(pl.gca(), make_mosaic(W, 6, 6), cmap=cm.binary)
```

USING THE SAME IDEA, ONE CAN ALSO VISUALIZE
THE FEATURE MAPS AT INTERMEDIATE LAYERS

THIS HELPS TRACING THE FEATURES LEARNED BY THE
NETWORK



KERAS:

```
def display_all_layer_filters(model, layer_dict, layer_name, input_img, dir_out=None):
    """
    if len(input_img.shape) == 2: # grey level image
        img = input_img.reshape(1, 1, input_img.shape[0], input_img.shape[1])
    elif len(input_img.shape) == 3:
        img = input_img.reshape(1, input_img.shape[0], input_img.shape[1], input_img.shape[2])
    else:
        raise TypeError("Input image should have 2 or 3 dimensions")

    get_layer_output = K.function([model.layers[0].input, K.learning_phase()],
[layer_dict[layer_name].output, K.learning_phase()])
    learning_ph = 0 # we are testing
    layer_output = get_layer_output([img, learning_ph])[0]

    number_of_filters = layer_output.shape[1]
    rows = np.uint8(np.floor(np.sqrt(number_of_filters)))
    cols = number_of_filters // rows
    rest = number_of_filters % rows
    if rest > 0:
        rows += 1
    fig, axarray = plt.subplots(rows, cols, sharex='col', sharey='row', figsize=(18,12))
    fig.suptitle(layer_name)

    for r in range(rows):
        for c in range(cols):
            index = r*rows*c
            if index >= number_of_filters:
                break
            # axarray[r,c].set_title()
            axarray[r,c].imshow(layer_output[0, index, :, :], interpolation="nearest", cmap="gray")
```

KERAS:

```
def display_all_layer_filters(model, layer_dict, layer_name, input_img, dir_out=None):
    """
    if len(input_img.shape) == 2: # grey level image
        img = input_img.reshape(1, 1, input_img.shape[0], input_img.shape[1])
    elif len(input_img.shape) == 3:
        img = input_img.reshape(1, input_img.shape[0], input_img.shape[1], input_img.shape[2])
    else:
        raise TypeError("Input image should have 2 or 3 dimensions")

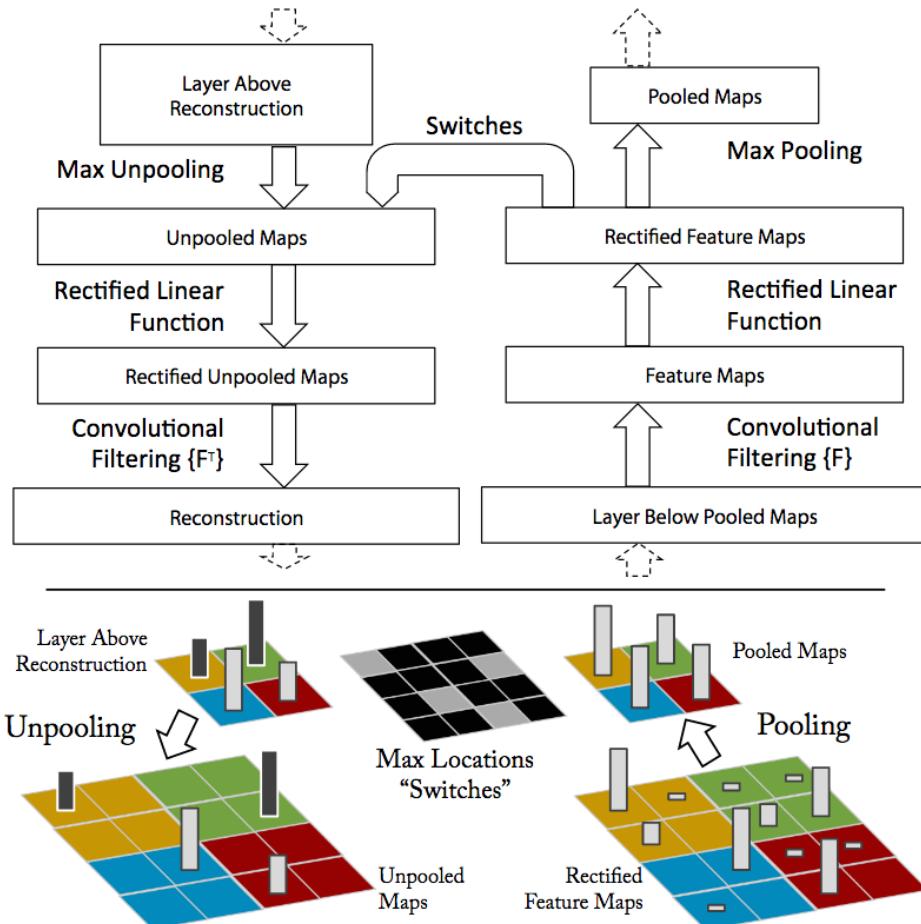
    get_layer_output = K.function([model.layers[0].input, K.learning_phase()],
        [layer_dict[layer_name].output, K.learning_phase()])
    learning_ph = 0 # we are testing
    layer_output = get_layer_output([img, learning_ph])[0]

    number_of_filters = layer_output.shape[1]
    rows = np.uint8(np.floor(np.sqrt(number_of_filters)))
    cols = number_of_filters // rows
    rest = number_of_filters % rows
    if rest > 0:
        rows += 1
    fig, axarray = plt.subplots(rows, cols, sharex='col', sharey='row', figsize=(18,12))
    fig.suptitle(layer_name)

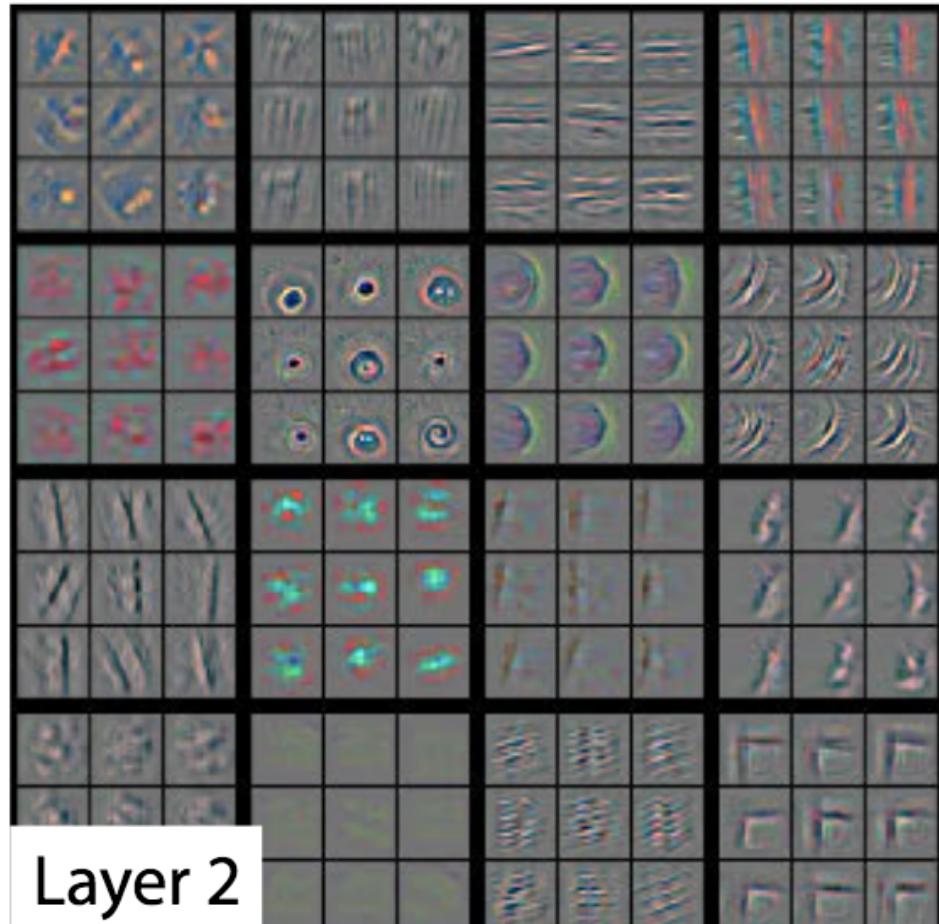
    for r in range(rows):
        for c in range(cols):
            index = r*rows*c
            if index >= number_of_filters:
                break
            # axarray[r,c].set_title()
            axarray[r,c].imshow(layer_output[0, index, :, :], interpolation="nearest", cmap="gray")
```

get layer output
for a series of “layer_names”

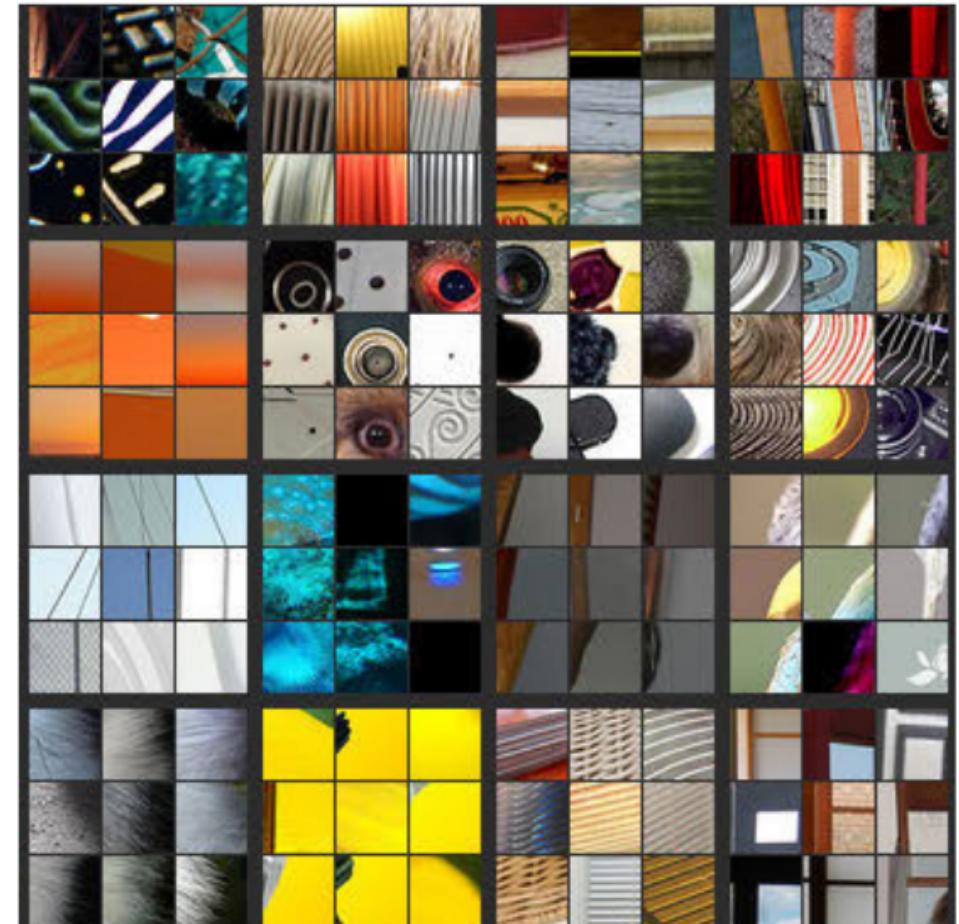
USE “DECONVNETS” TO MAP BACK THE FEATURE MAP INTO THE PIXEL SPACE



IT ALLOWS TO SEE
WHICH
REGIONS OF THE INPUT
GENERATED
A MAXIMUM RESPONSE
IN A NEURON

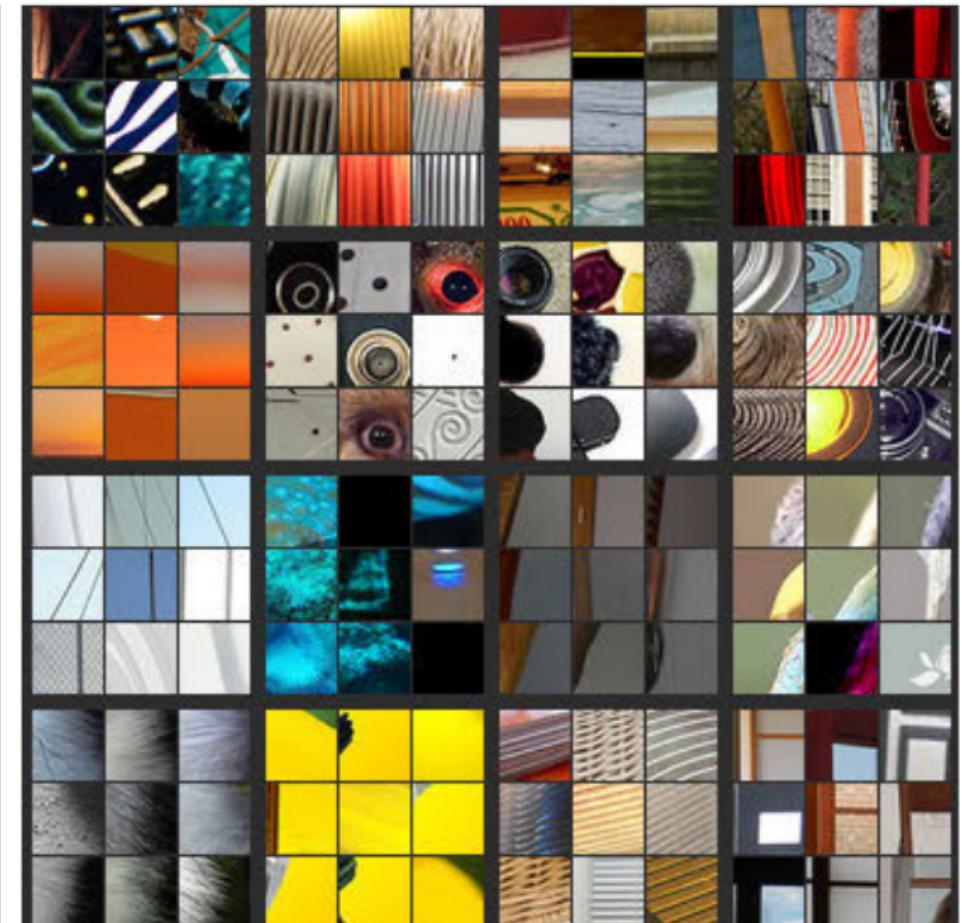
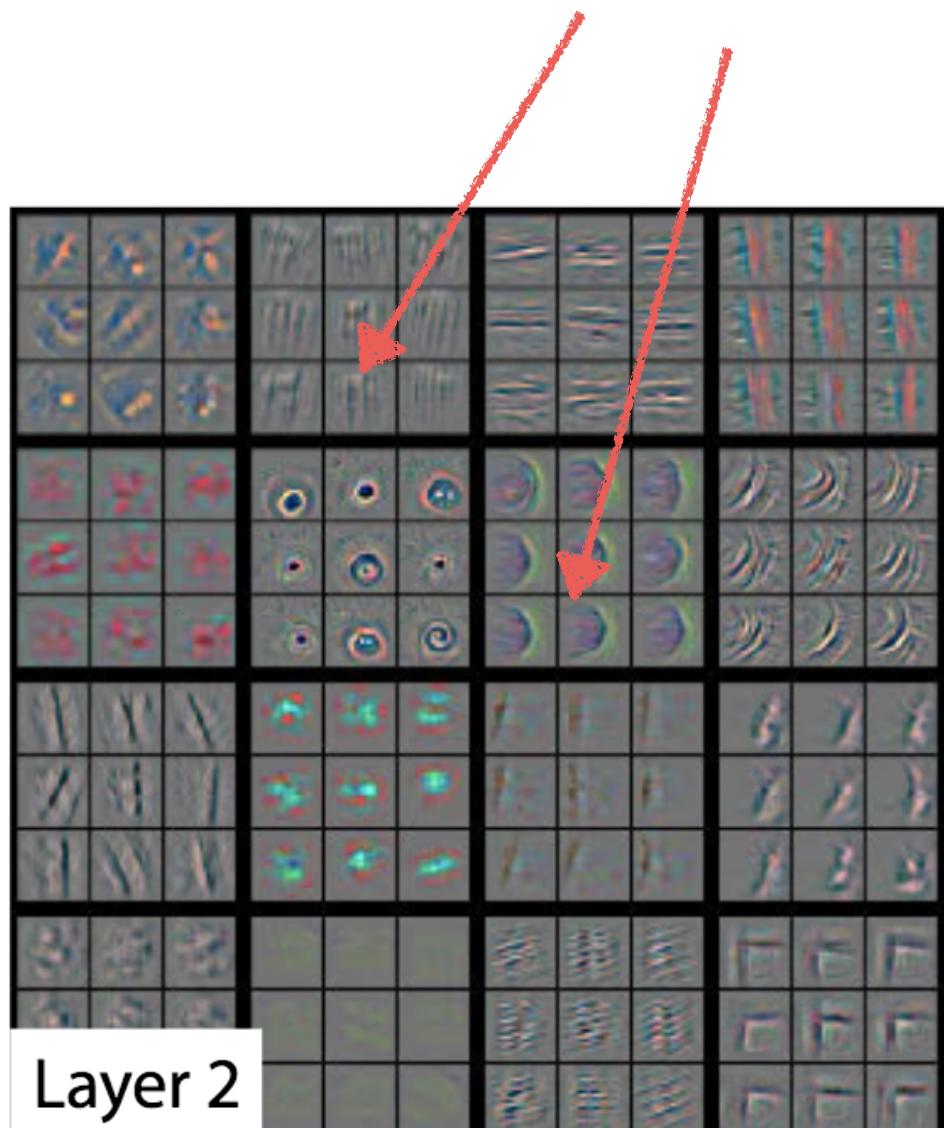


Layer 2

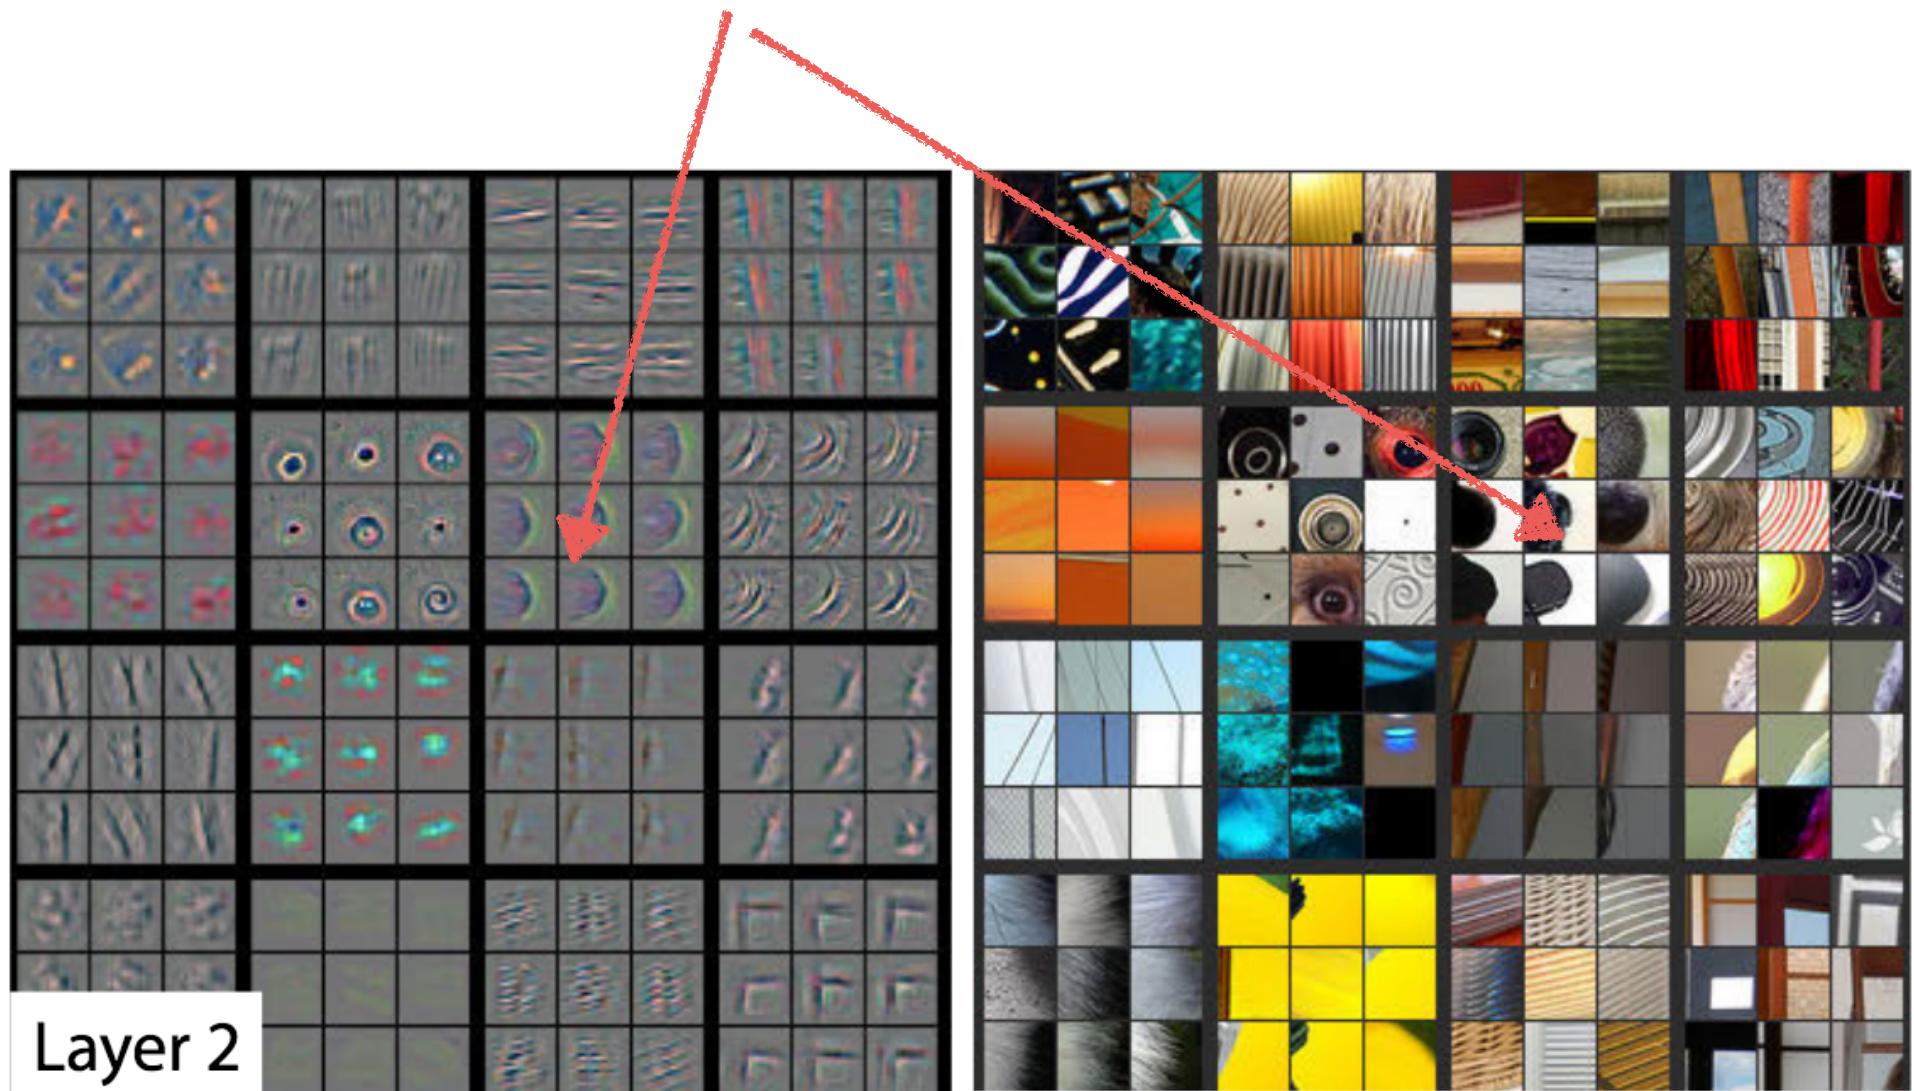


Zeiler+14

EVERY BLOCK OF 9 SHOWS
THE 9 STRONGEST RESPONSES TO A GIVEN FILTER OF LAYER2



THE CORRESPONDING REGIONS OF IMAGES THAT GENERATED THE MAXIMUM RESPONSE

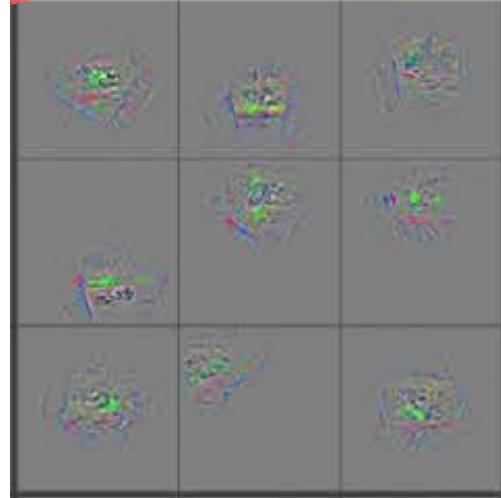


CAN BE
REPEATED
FOR DEEPER
LAYERS
ALTHOUGH IT
BECOMES LESS
INTUITIVE

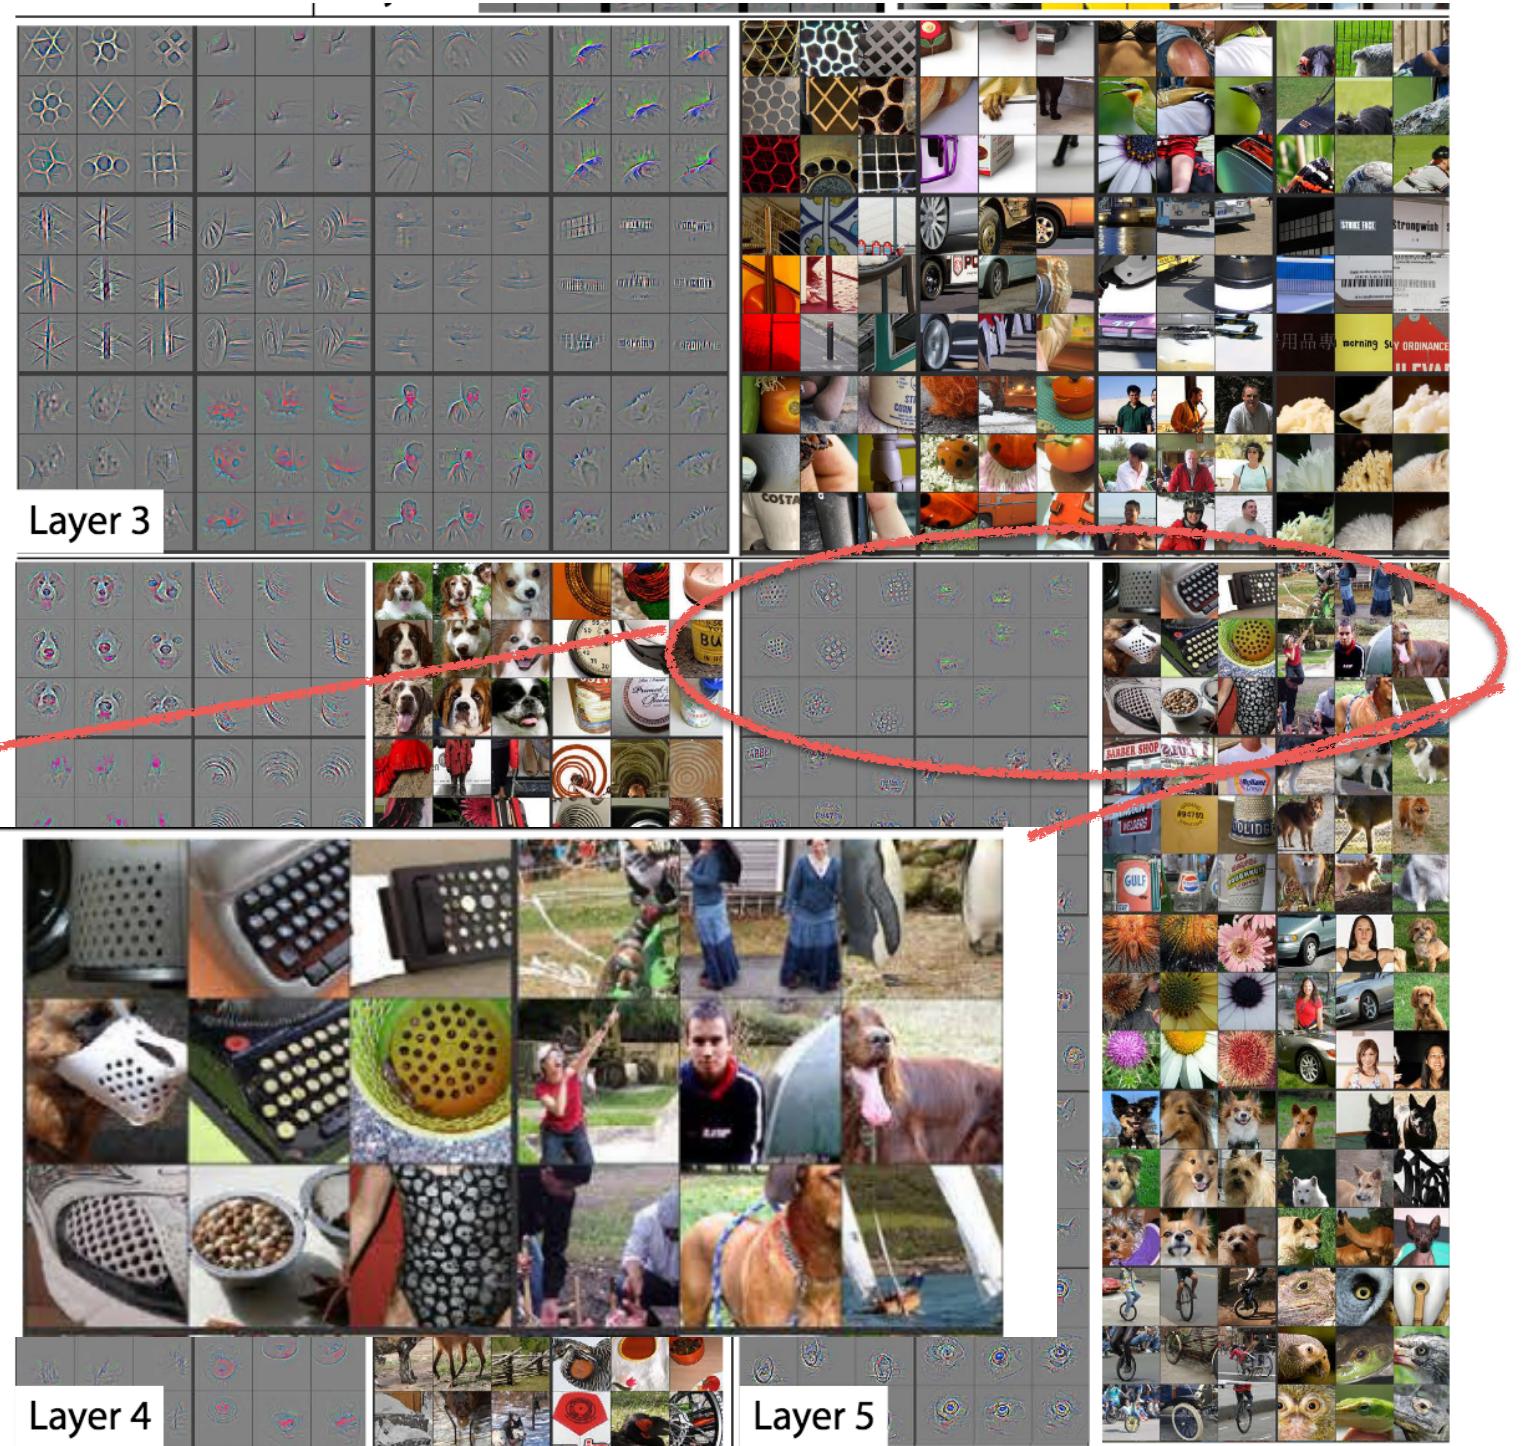
Zeiler+14



CAN BE
REPEATED
FOR DEEPER
LAYERS
ALTHOUGH IT
BECOMES LESS



Zeiler+14



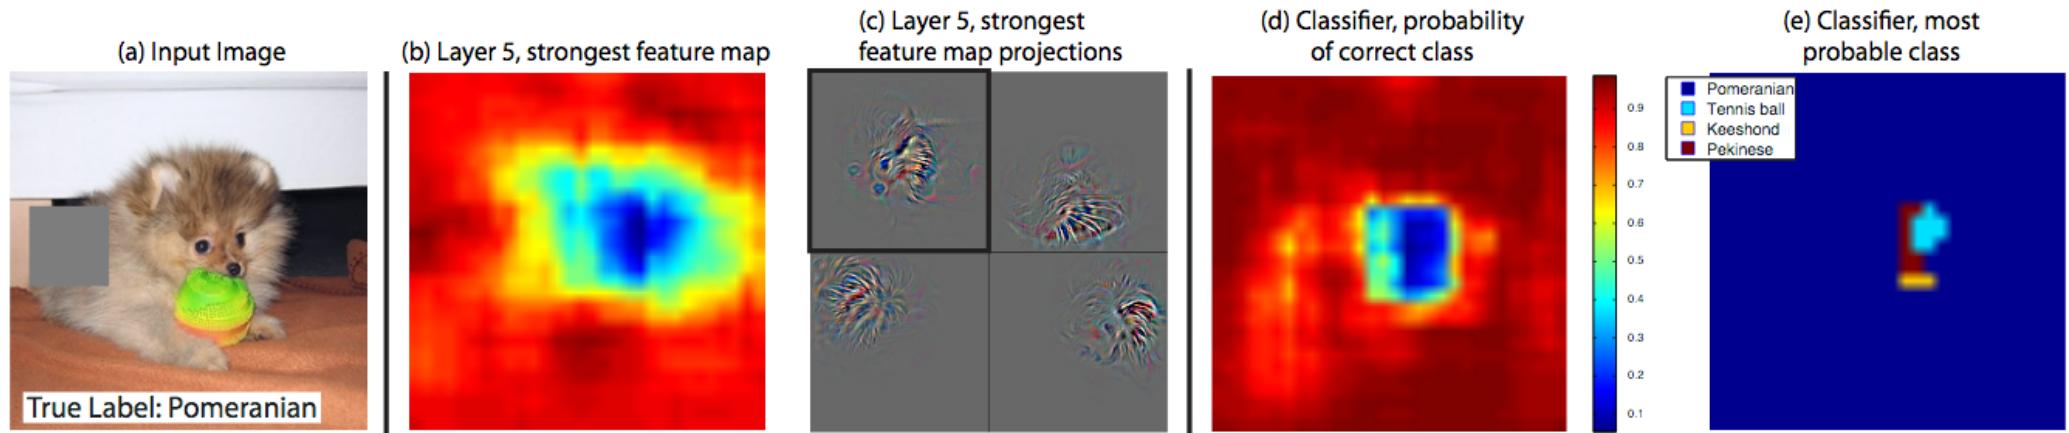
KERAS IMPLEMENTATION OF VISUALIZATIONS THROUGH DECONVNETS

<https://github.com/jalused/Deconvnet-keras>

OCCLUSION SENSITIVITY TRIES ALSO TO FIND THE REGION OF THE IMAGE THAT TRIGGERED THE NETWORK DECISION BY MASKING DIFFERENT REGIONS OF THE INPUT IMAGE AND ANALYZING THE NETWORK OUTPUT

IT ALLOWS TO IF THE NETWORK IS TAKING THE DECISIONS BASED ON THE EXPECTED FEATURES

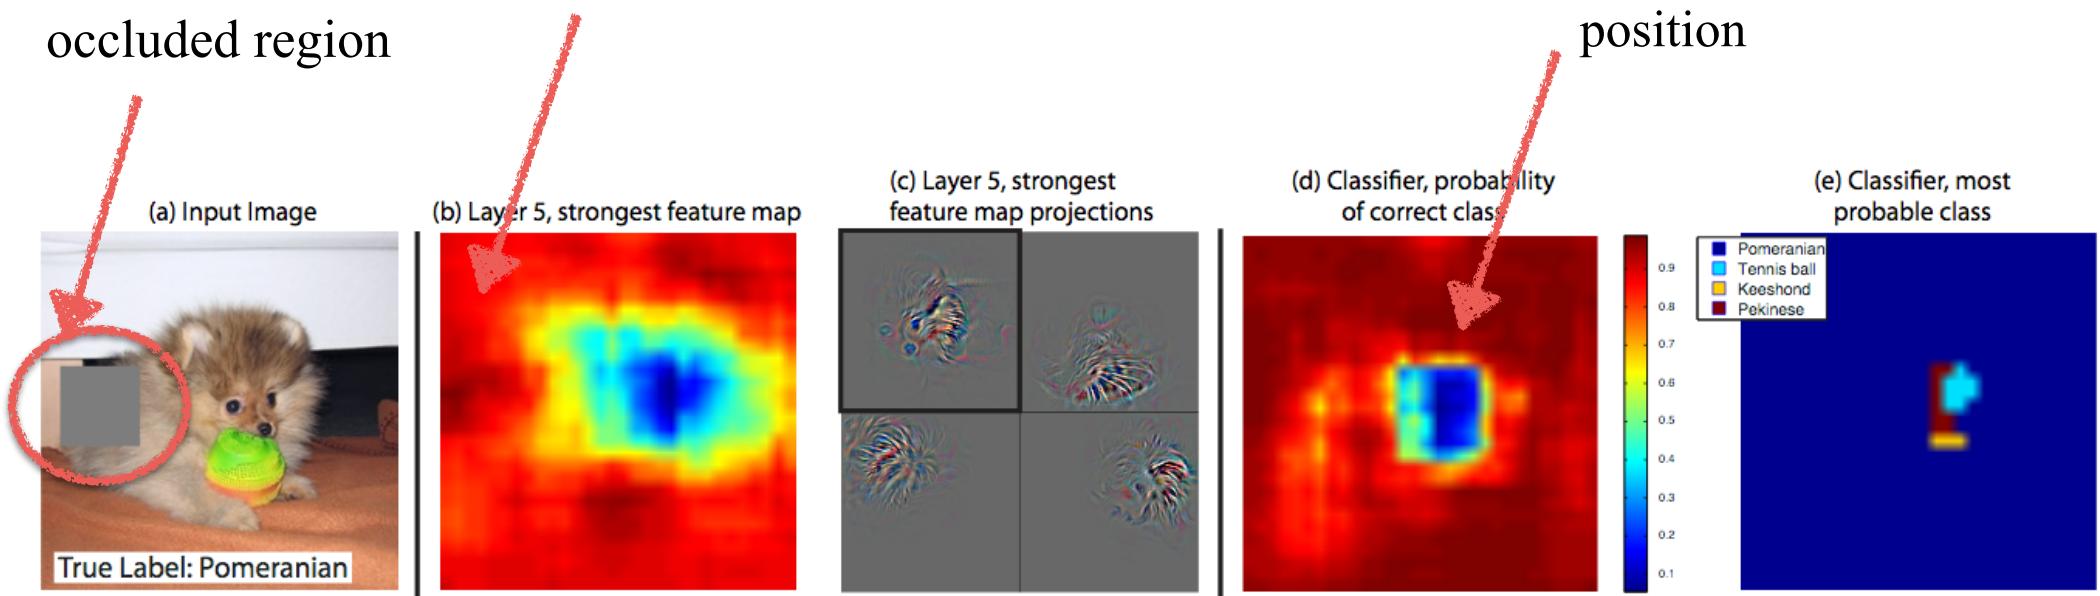
VERY TIME CONSUMING!



OCCLUSION SENSITIVITY TRIES ALSO TO FIND THE REGION OF THE IMAGE THAT TRIGGERED THE NETWORK DECISION BY MASKING DIFFERENT REGIONS OF THE INPUT IMAGE AND ANALYZING THE NETWORK OUTPUT

for every position
of the square the maximum response of a given layer
is averaged

occluded region



INCEPTIONISM - DEEP DREAM

THE IDEA BEHIND INCEPTIONISM TECHNIQUES
IS TO INVERT THE NETWORK TO GENERATE AN IMAGE
THAT MAXIMIZES THE OUTPUT SCORE

$$\arg \max_I S_c(I) - \lambda ||I||_2^2$$

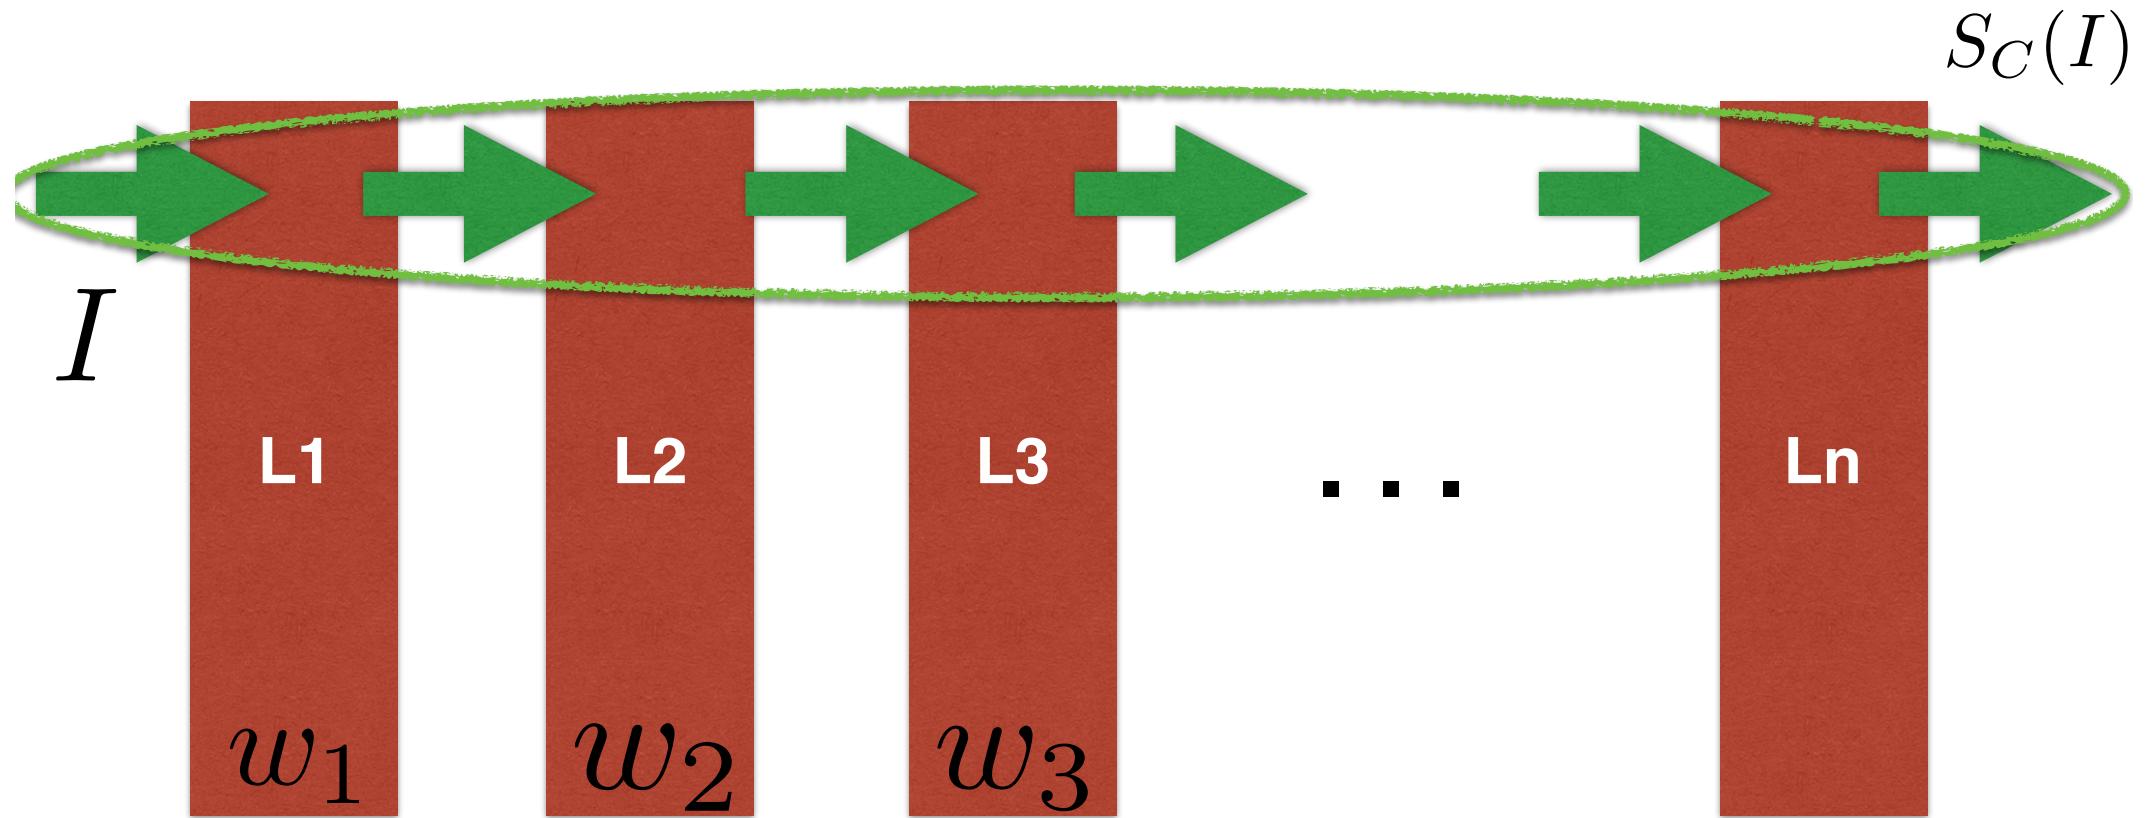
Score of class c for image I

image I

The diagram illustrates the mathematical expression for generating an image that maximizes the output score for a specific class. The expression is $\arg \max_I S_c(I) - \lambda ||I||_2^2$. Two red arrows point to the terms: one from the text 'Score of class c for image I' to the term $S_c(I)$, and another from the text 'image I' to the term $||I||_2^2$.

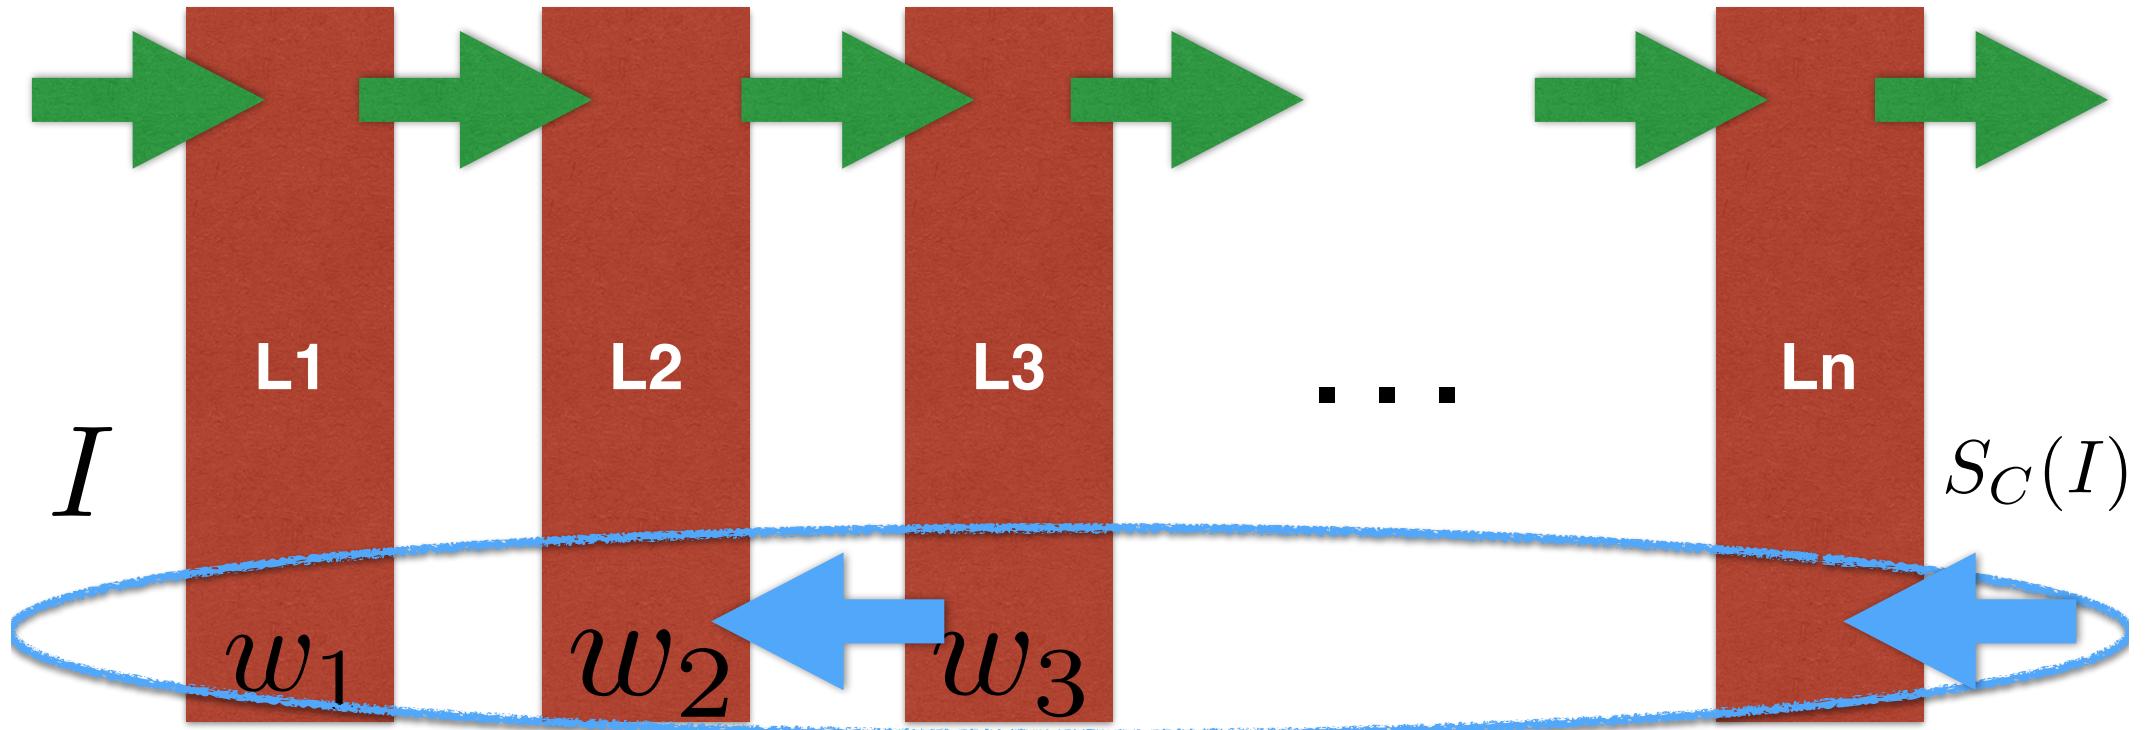
TRY TO FIND AN IMAGE THAT GENERATES A
HIGH SCORE FOR A GIVEN CLASS

INCEPTIONISM - DEEP DREAM



DURING THE TRAINING PHASE THE WEIGHTS ARE
LEARNED TO MAP I INTO S_C

INCEPTIONISM - DEEP DREAM



DURING THE RECONSTRUCTION PHASE, I IS LEARNT
THROUGH BACKPROPAGATION KEEPING THE WEIGHTS
FIXED

INCEPTIONISM - DEEP DREAM

RESULTS REVEAL INTERESTING INFORMATION ON
HOW THE NETWORKS BUILD REPRESENTATIONS OF
OBJECTS



Hartebeest



Measuring Cup



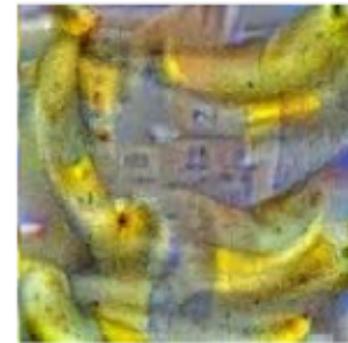
Ant



Starfish



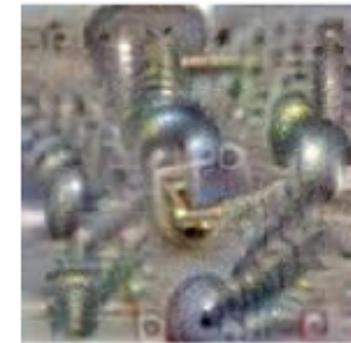
Anemone Fish



Banana



Parachute



Screw

INCEPTIONISM - DEEP DREAM

RESULTS REVEAL INTERESTING INFORMATION ON
HOW THE NETWORKS BUILD REPRESENTATIONS OF
OBJECTS



SOME STRANGE CASES...

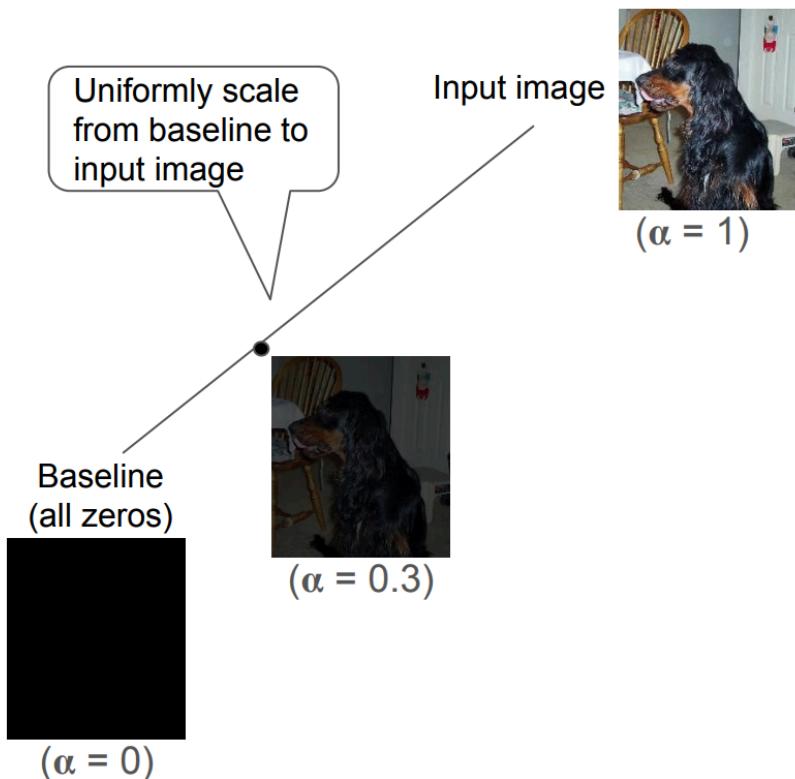
DEEP DREAM

<https://deeppdreamgenerator.com/>

IT HAS NOW BECOME A SORT OF ART?



INTEGRATED GRADIENTS



Construct a sequence of images interpolating from a baseline (black) to the actual image

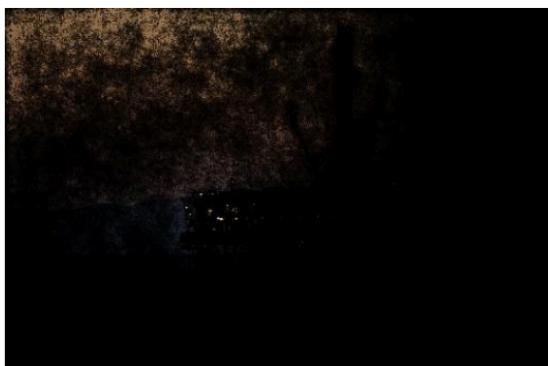
Average the gradients across these images

INTEGRATED GRADIENTS

Original image (Drilling platform)



Gradient at image



Integrated gradient

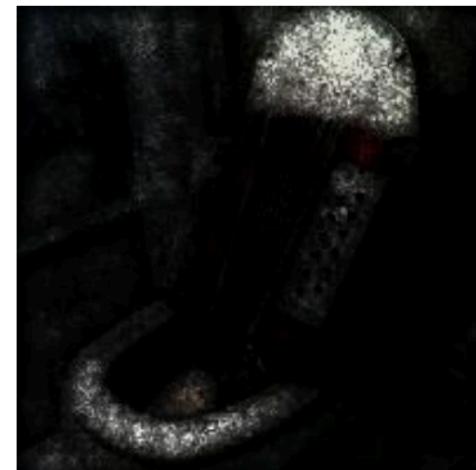


INTEGRATED GRADIENTS

Human label: **accordion**
Network's top label: **toaster**

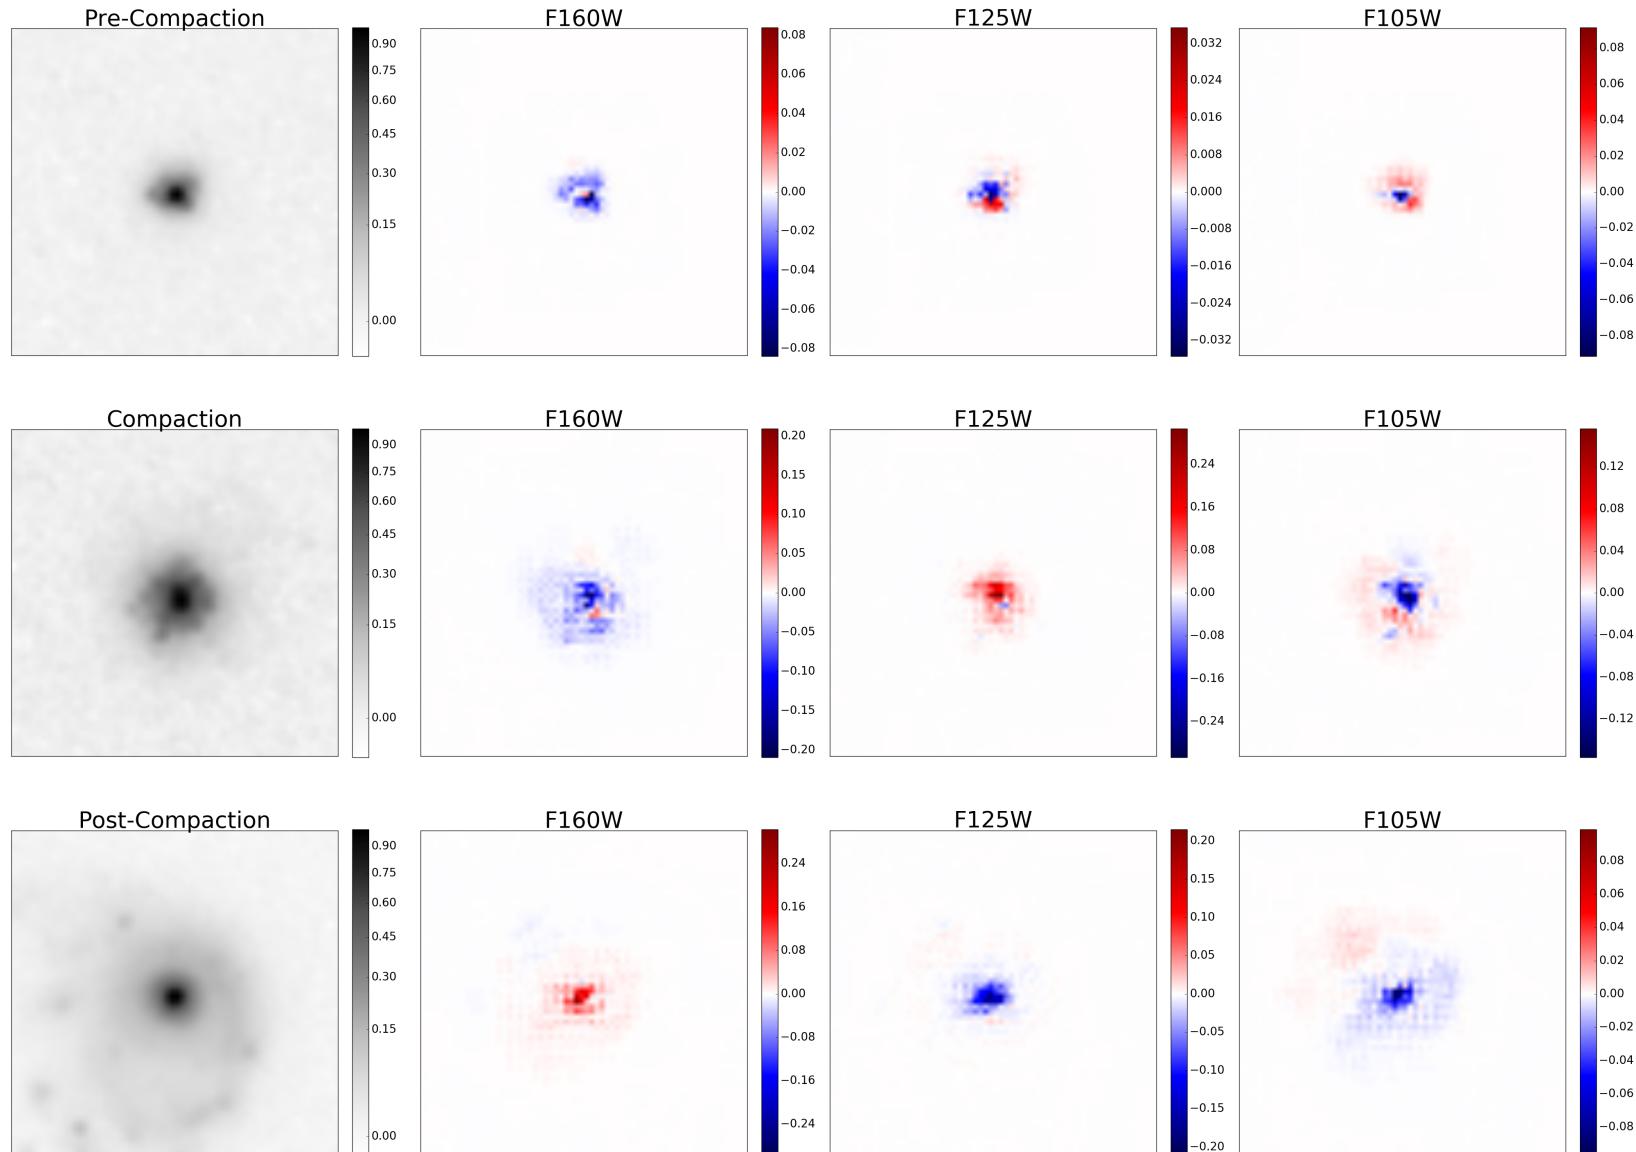


Integrated gradient



INTEGRATED GRADIENTS

[Applied to galaxies]



INTEGRATED GRADIENTS

KERAS IMPLEMENTATION:

<https://github.com/hiranumn/IntegratedGradients>